



© iStock / Fraunhofer ITWM

# CARME

An Open Source Framework for Multi-User,  
Interactive Jobs on Distributed GPU-Systems

The Carme Team  
Competence Center for High Performance Computing,  
Fraunhofer ITWM, Kaiserslautern, GERMANY



# Machine Learning in the Wild

Why we need a tool for ML/DL on HPC Clusters

# Machine Learning in the wild

## Why we need a tool for ML/DL on HPC Clusters

### Machine Learning and Data Analytics



showed **remarkable success** in many distinct fields  
(e.g. image processing, text translation, medical imaging, ... )



rise of **hybrid** and/or **ML/DL** supported **algorithms**  
(e.g. climate forecast, physics informed networks, ... )



**model size increases** drastically  
(needs more memory & compute power)

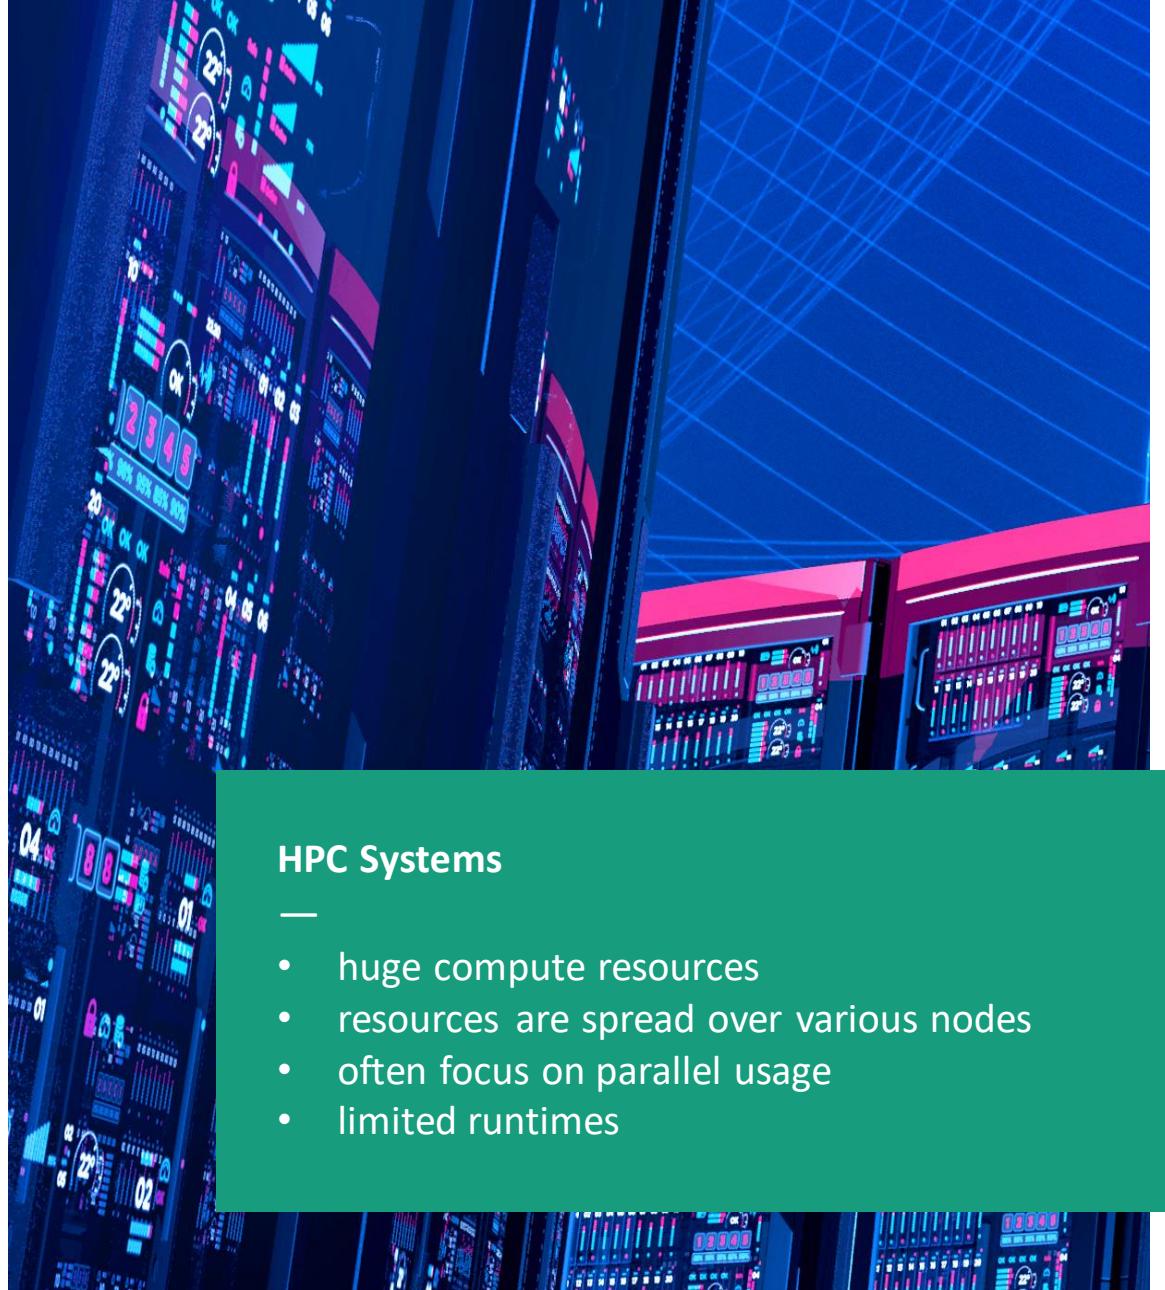


generate **large investments** in new **multi-GPU hardware**

# Utilize HPC Systems

The way it was for decades

- other fields use **HPC systems for decades**  
(e.g. physics, chemistry, mathematics, ... )
- HPC clusters are set up to **make use of the hardware** that is installed  
(e.g. multi-node jobs)
- **calculate huge problems**  
(e.g. extreme memory and computationally intensive simulations)
- **commandline driven** usage  
(e.g. submit non-interactive batch jobs, jobs do not start directly)



## HPC Systems

- huge compute resources
- resources are spread over various nodes
- often focus on parallel usage
- limited runtimes

# Machine Learning and HPC Systems

## What we would expect

### Machine Learning and HPC Systems



efficiently **manage and use your resources**  
(how to utilize the full power of multi-GPU systems?)



**match your workflow** with these resources  
(how can I develop interactively?)



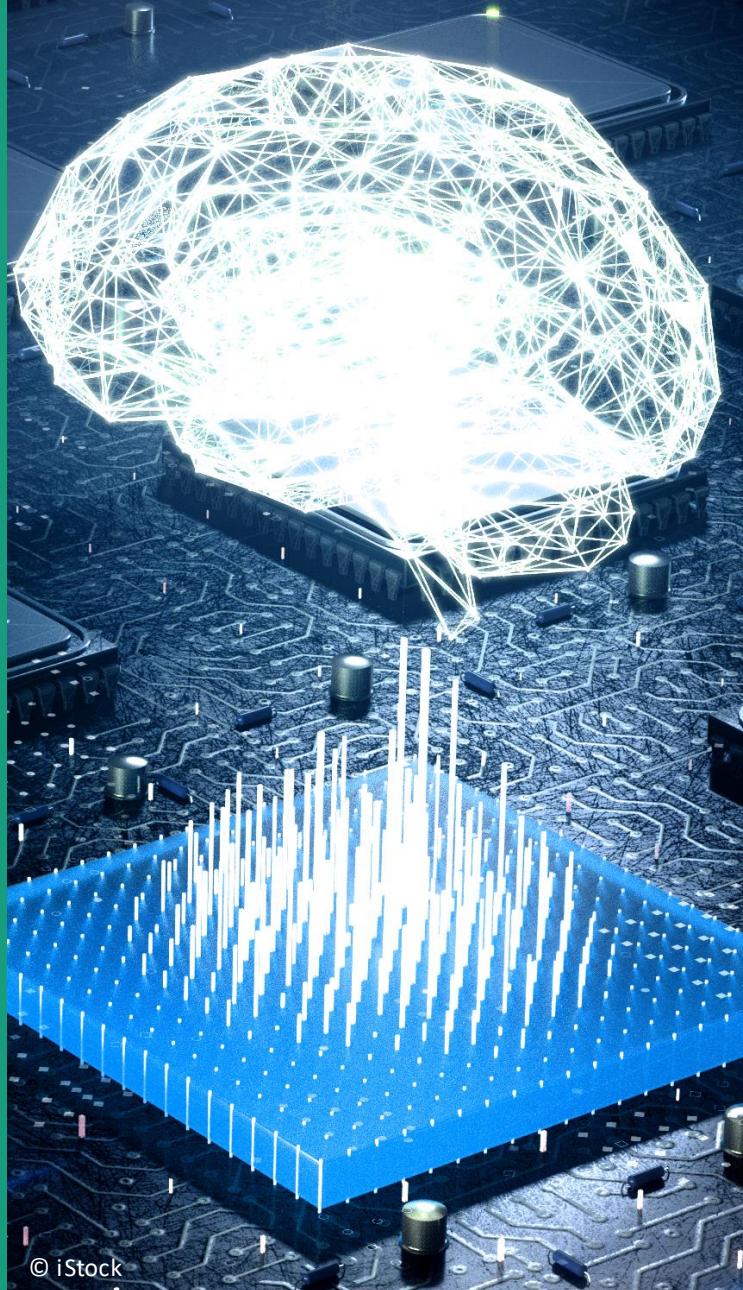
scale your applications to **use multiple GPUs and/or nodes**  
(how do I easily utilize multiple GPUs in various nodes?)

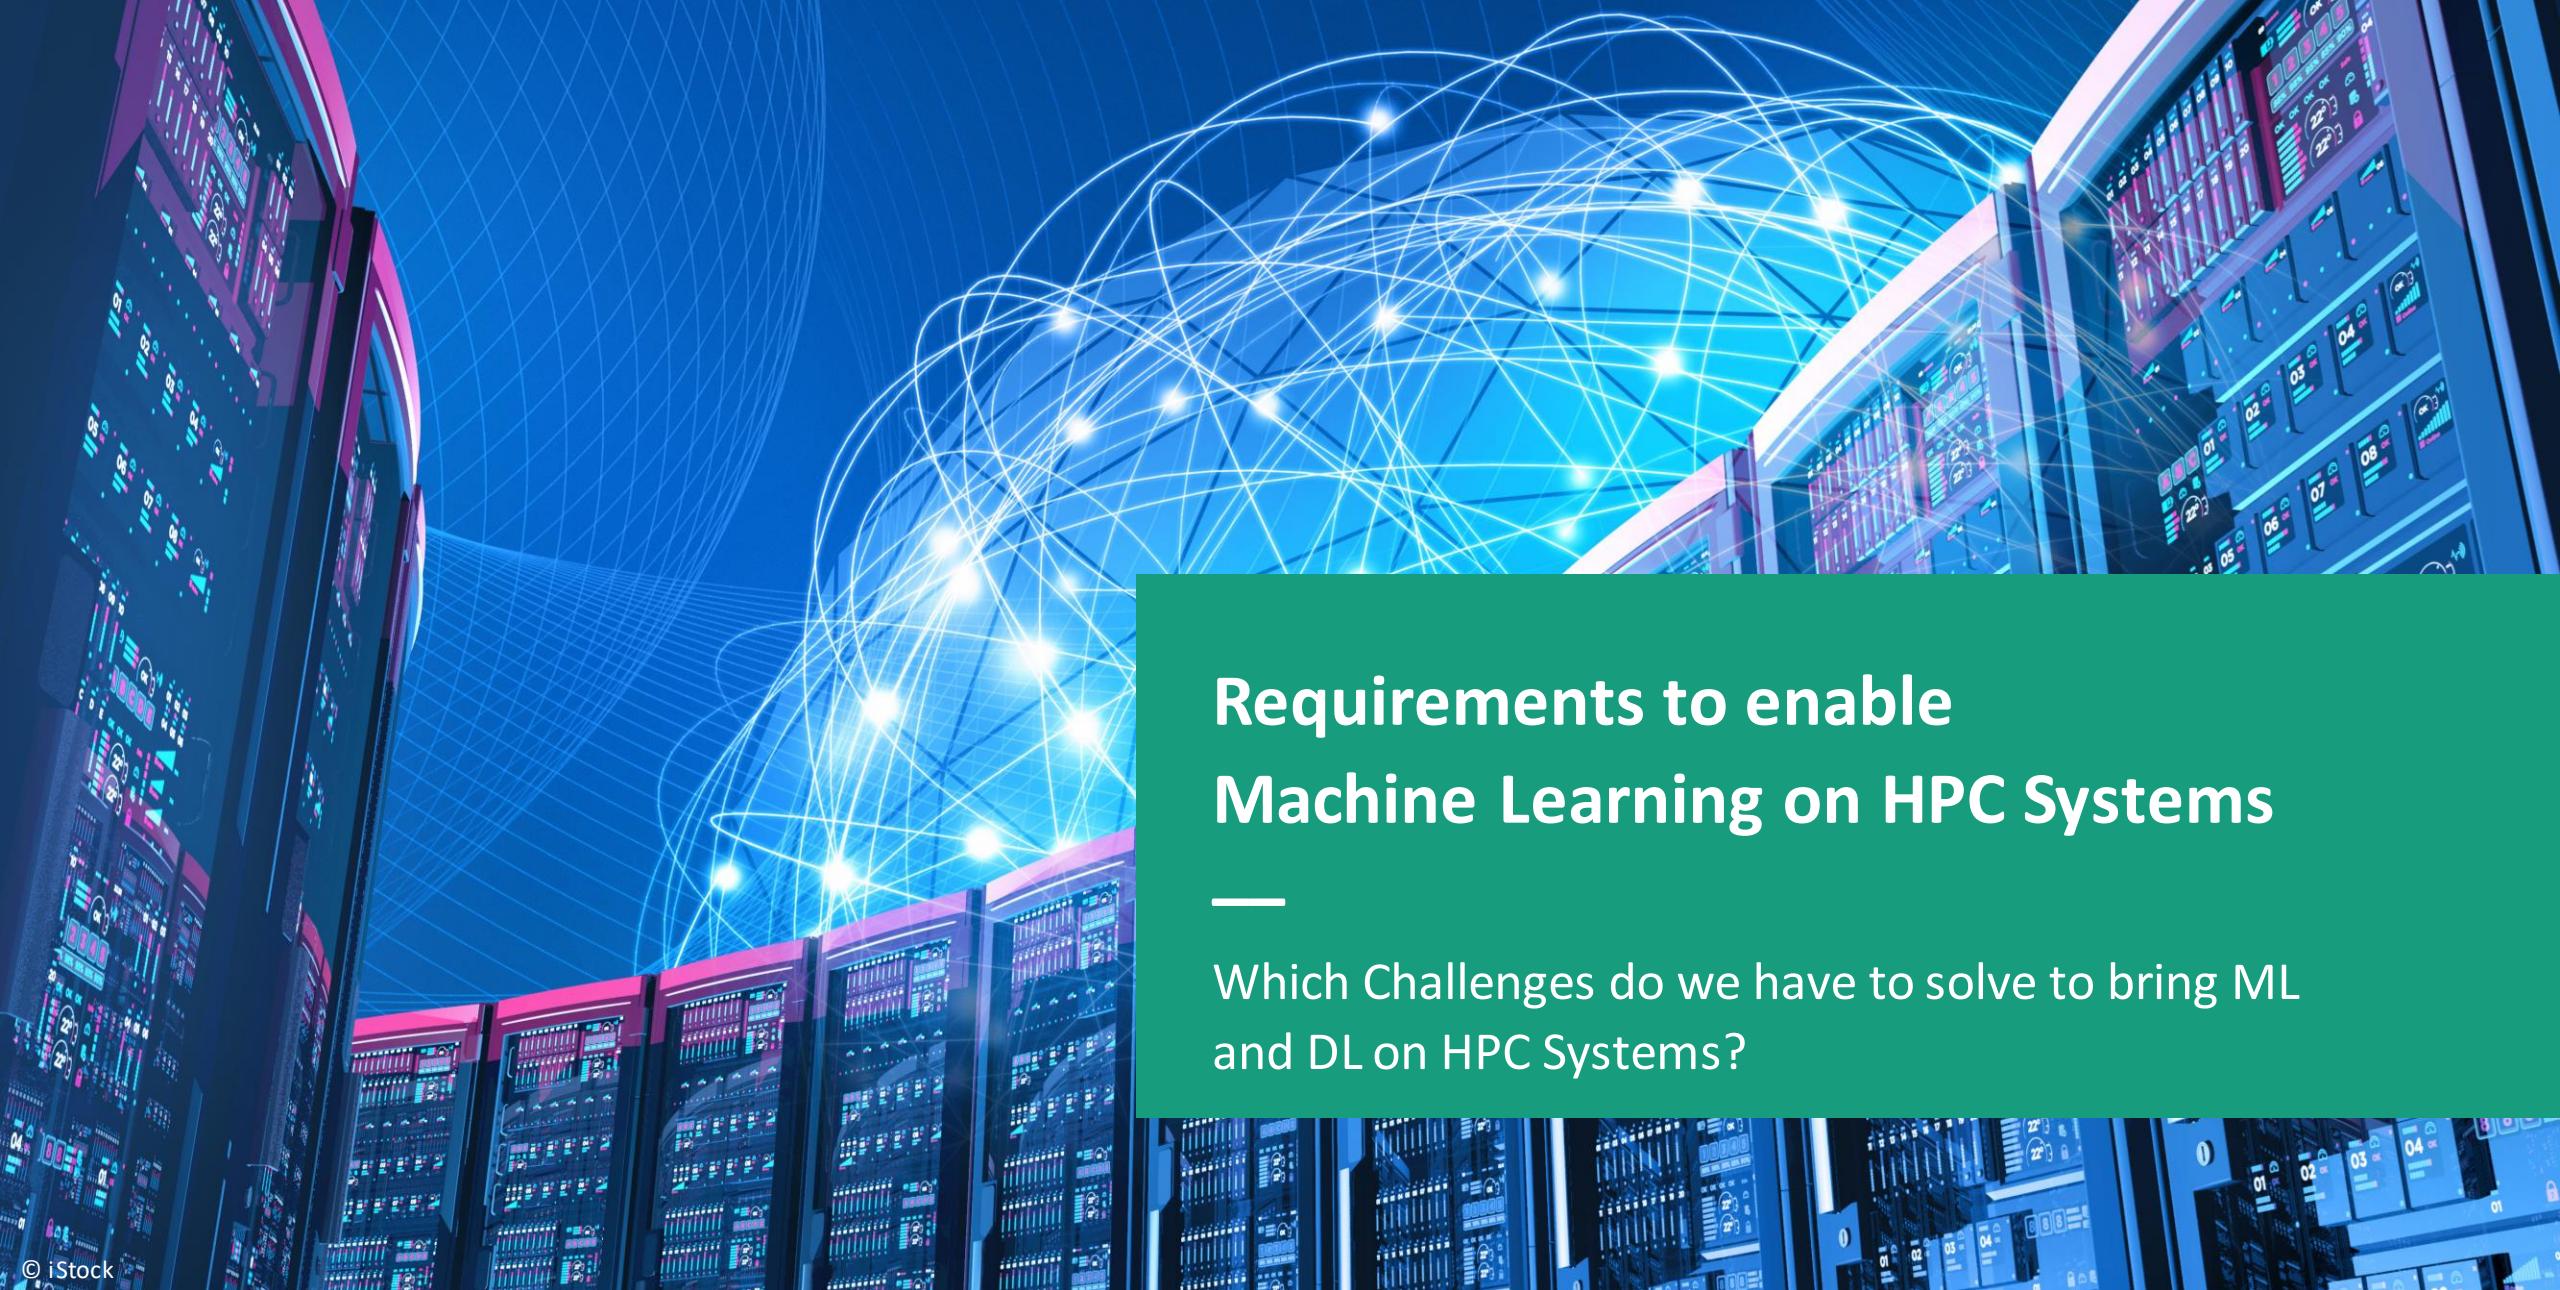


manage **data I/O and data storage**  
(how do I access my data?)



# Machine Learning and HPC Systems two distinct worlds?



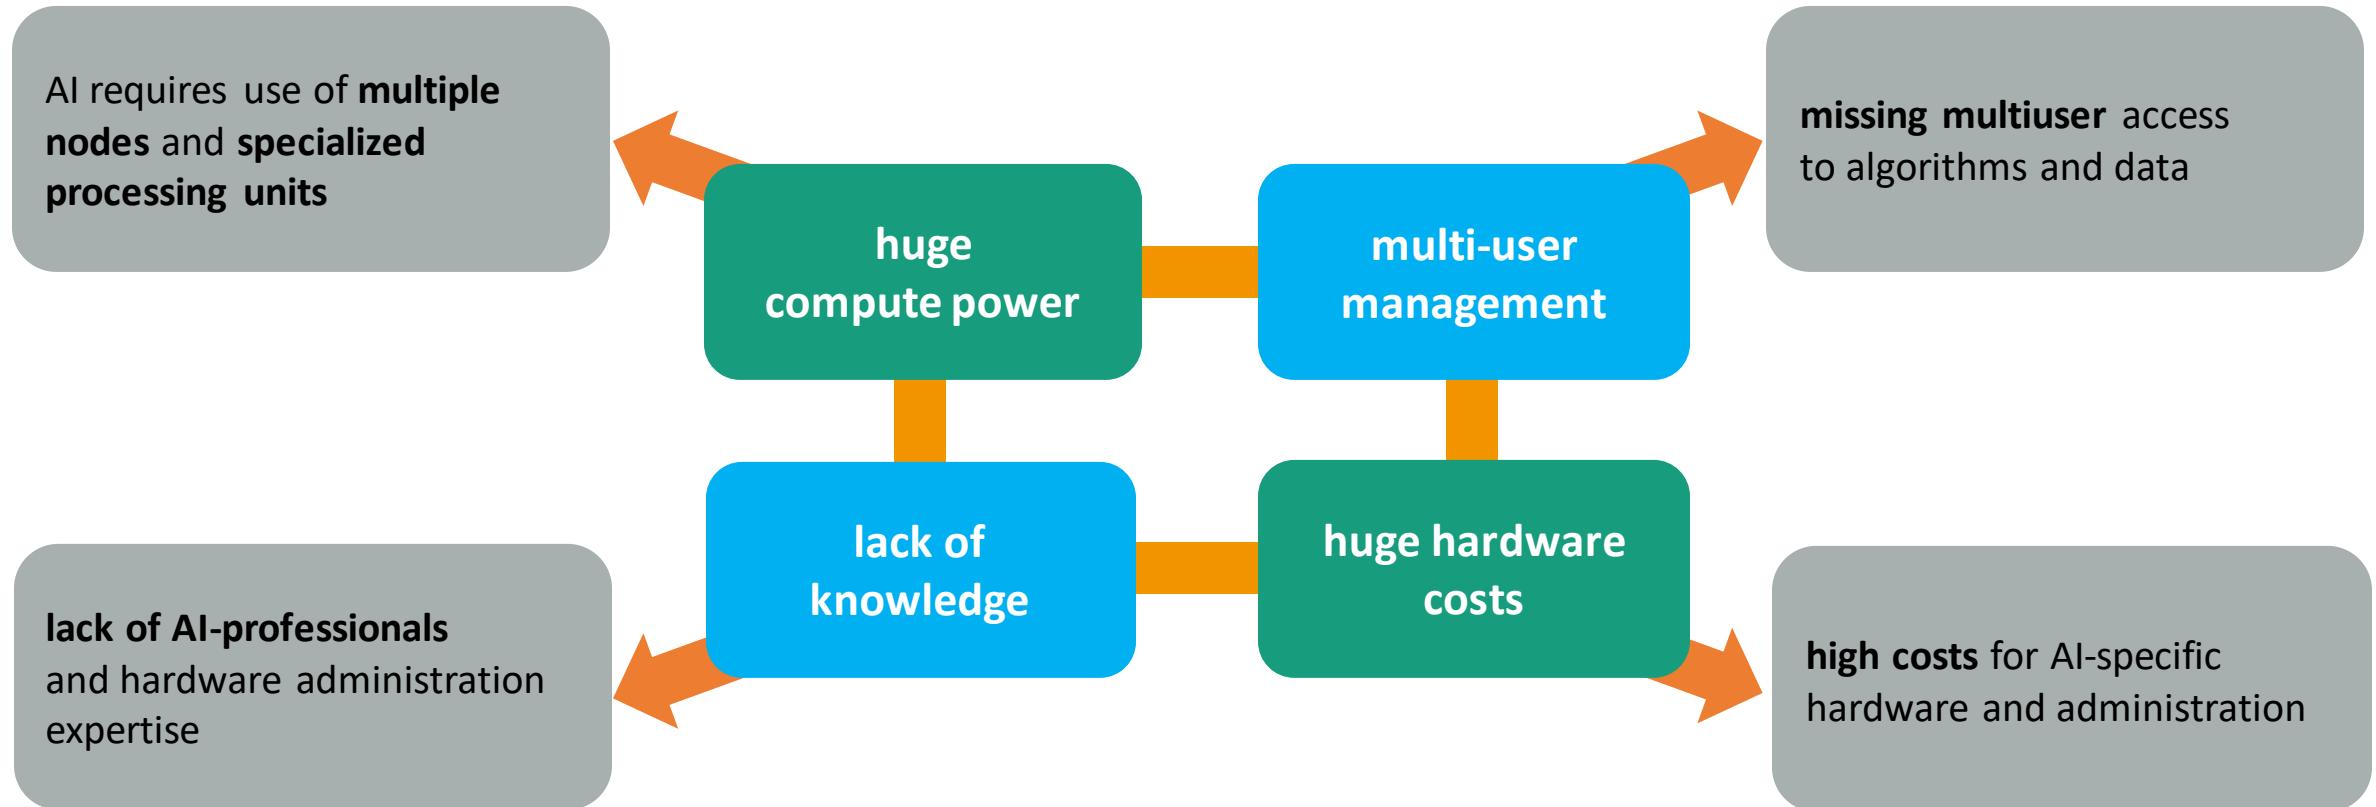


## Requirements to enable Machine Learning on HPC Systems

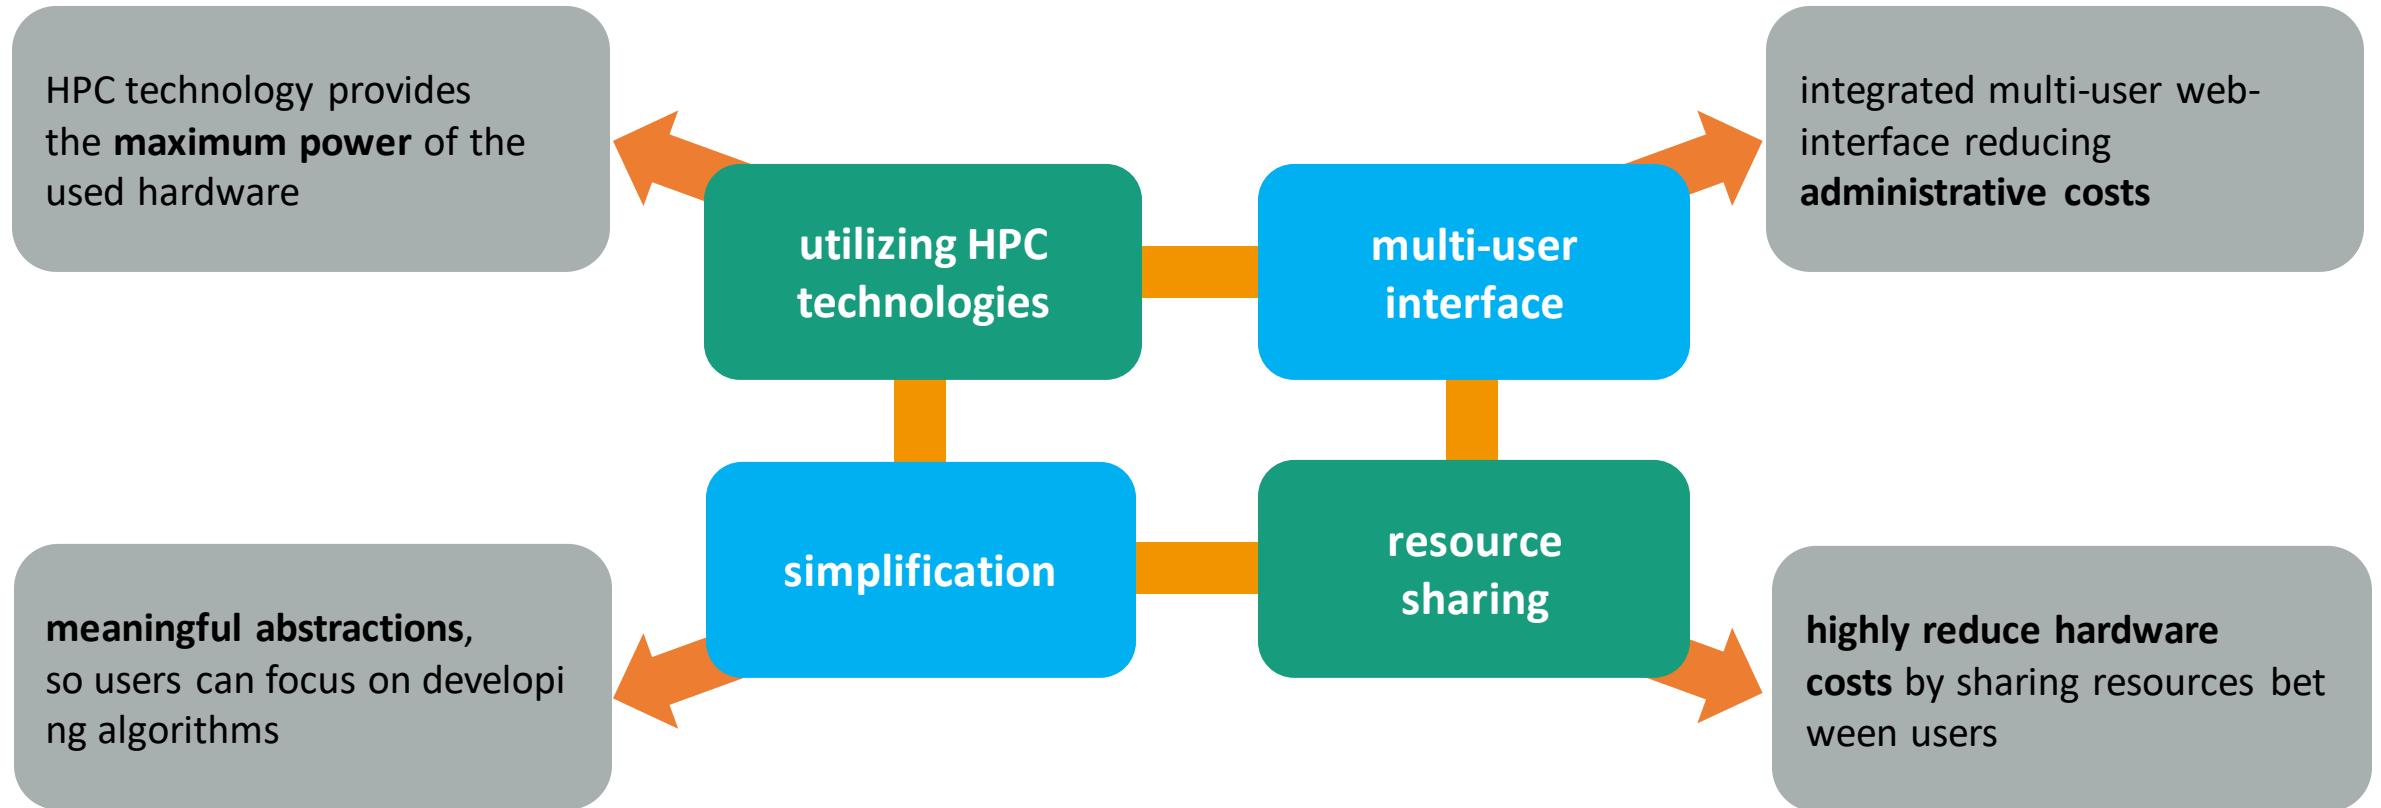
Which Challenges do we have to solve to bring ML and DL on HPC Systems?

# Machine Learning and HPC Systems

on an HPC system we have to take care of ...



## our proposed solutions ...

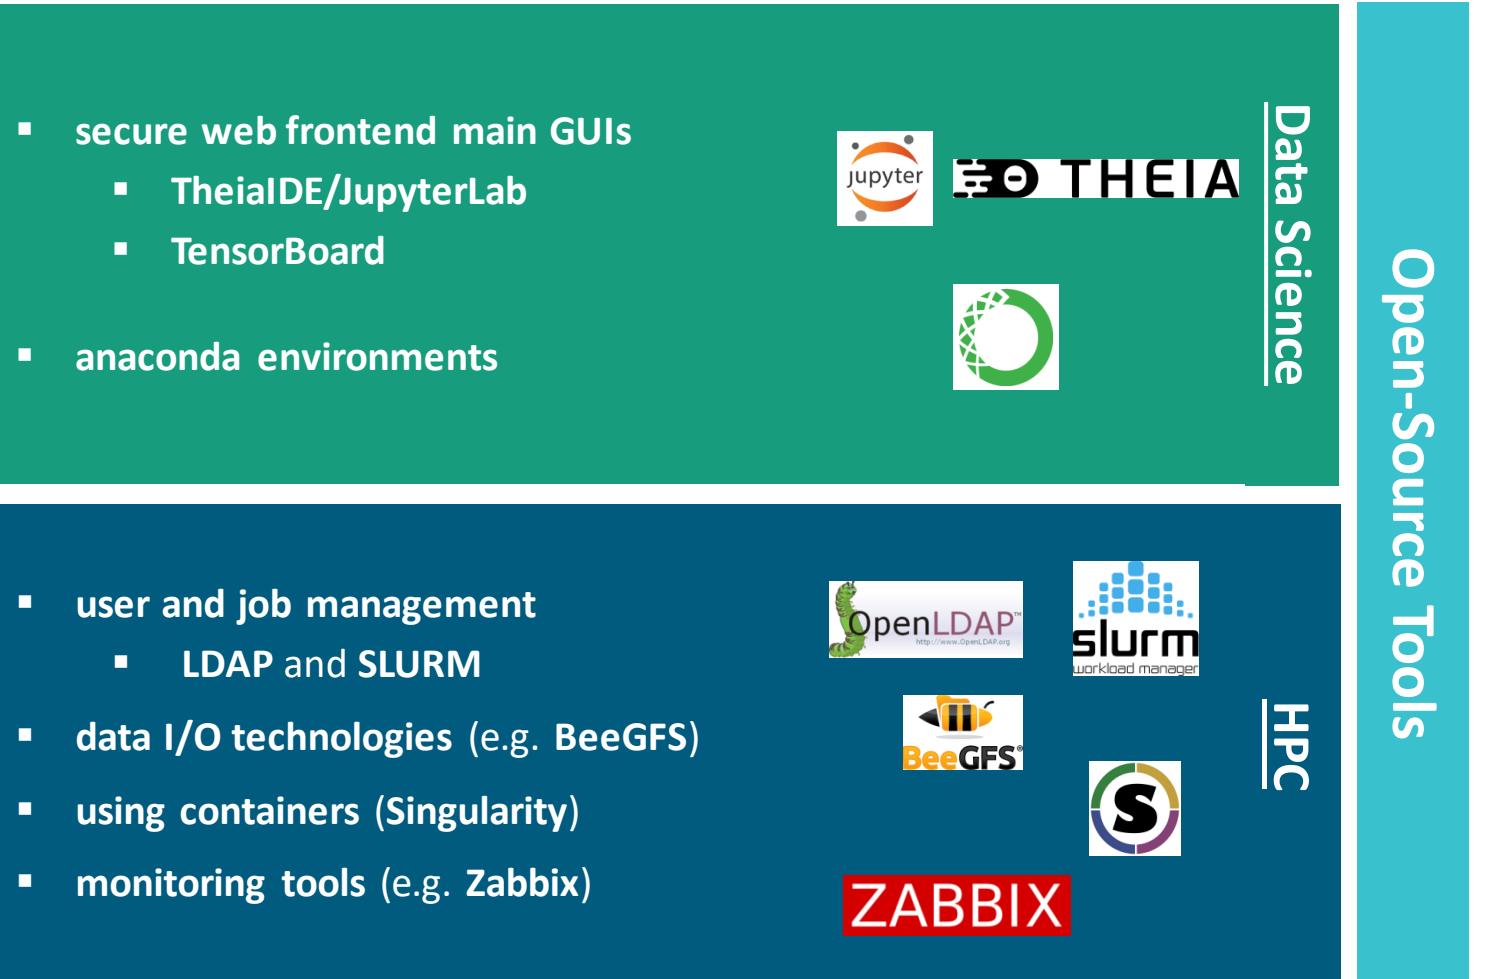


# Combining Open-Source ML Tools with HPC Backends

## The CARME Way



**Carme** – an open source framework to manage resources for multiple users running interactive jobs on an HPC Clusters



<https://carmeteam.github.io/Carme/>



## The Details of our Solution

—  
What it is really about



# Resource Management

How to assign resources to competing users?

# Resource Management

## Proven HPC Tools

### LDAP<sup>[1]</sup> (user management)



- LDAP = Lightweight Directory Access Protocol
- use **your LDAP**
- **different user roles**

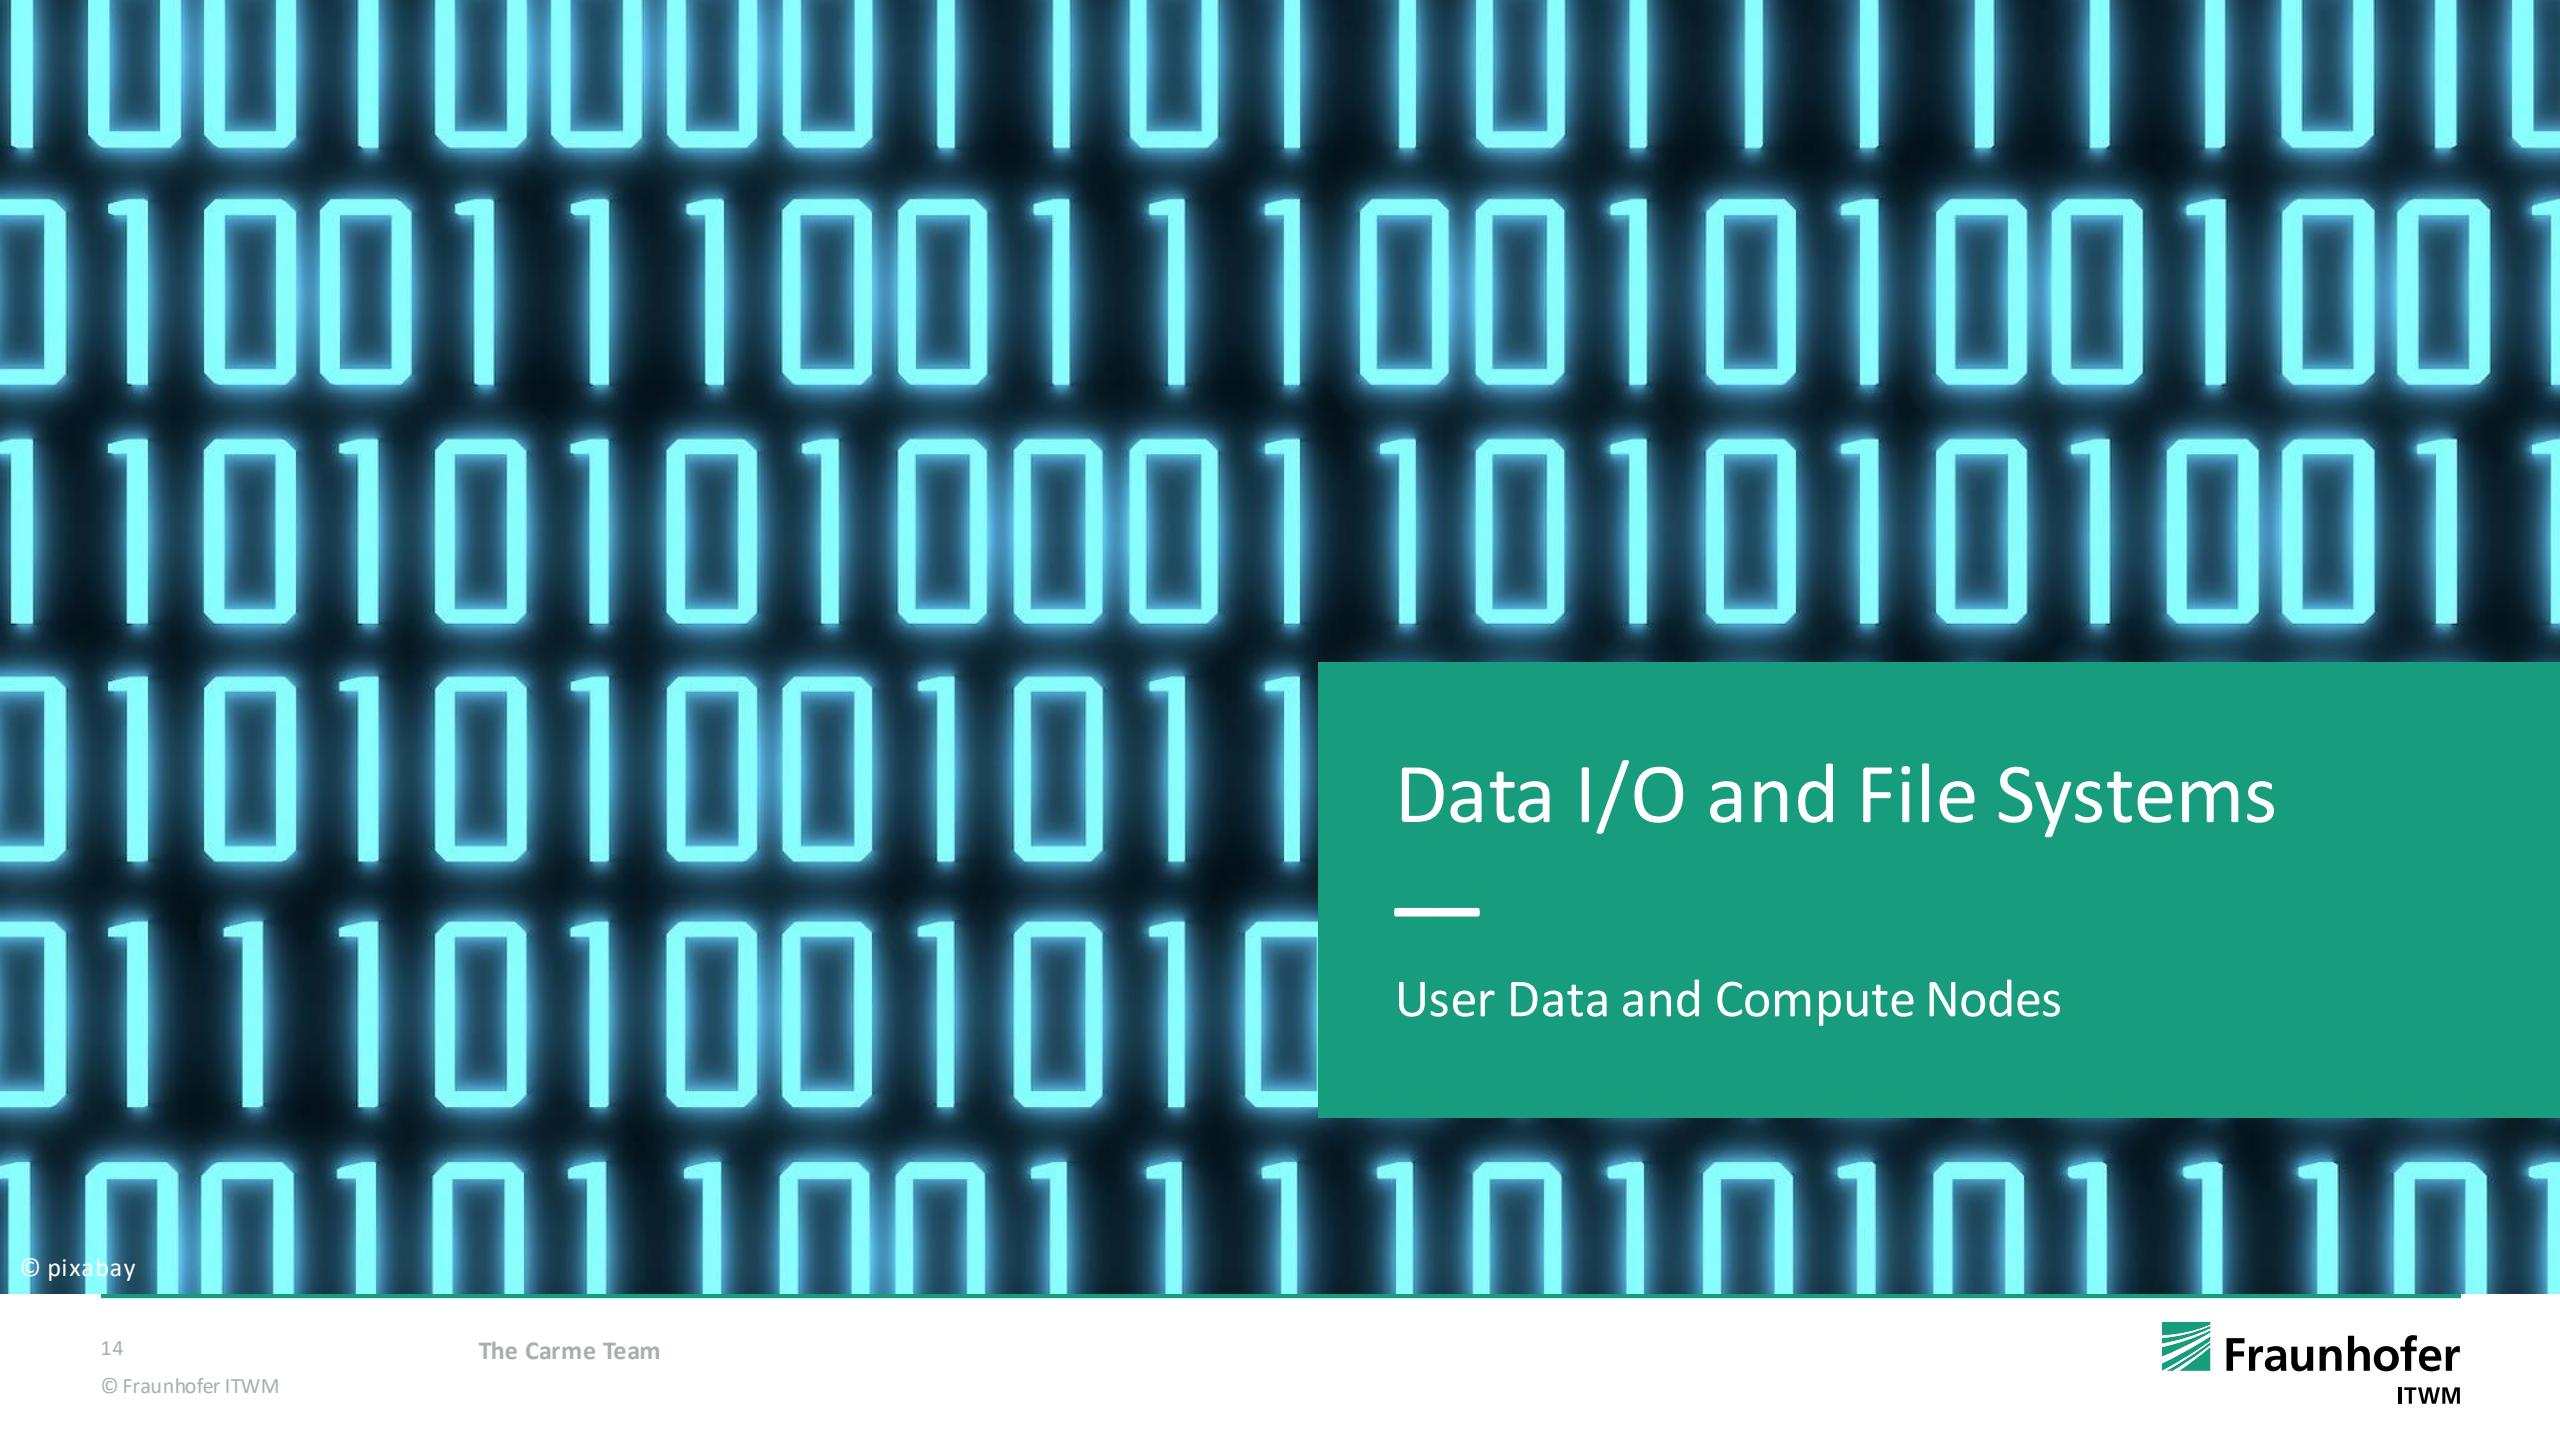
### SLURM<sup>[2]</sup> (job scheduler)



- easy **integration in existing installations**
- make use of **cgroups** plugins
- using **quotas, queues, resource reservation**
- many **useful extensions** available  
(e.g. preemption, hibernate if idle, ...)

[1] <https://www.openldap.org>

[2] <https://slurm.schedmd.com>



## Data I/O and File Systems

---

User Data and Compute Nodes

# Data I/O and File Systems

Distributed for more Power

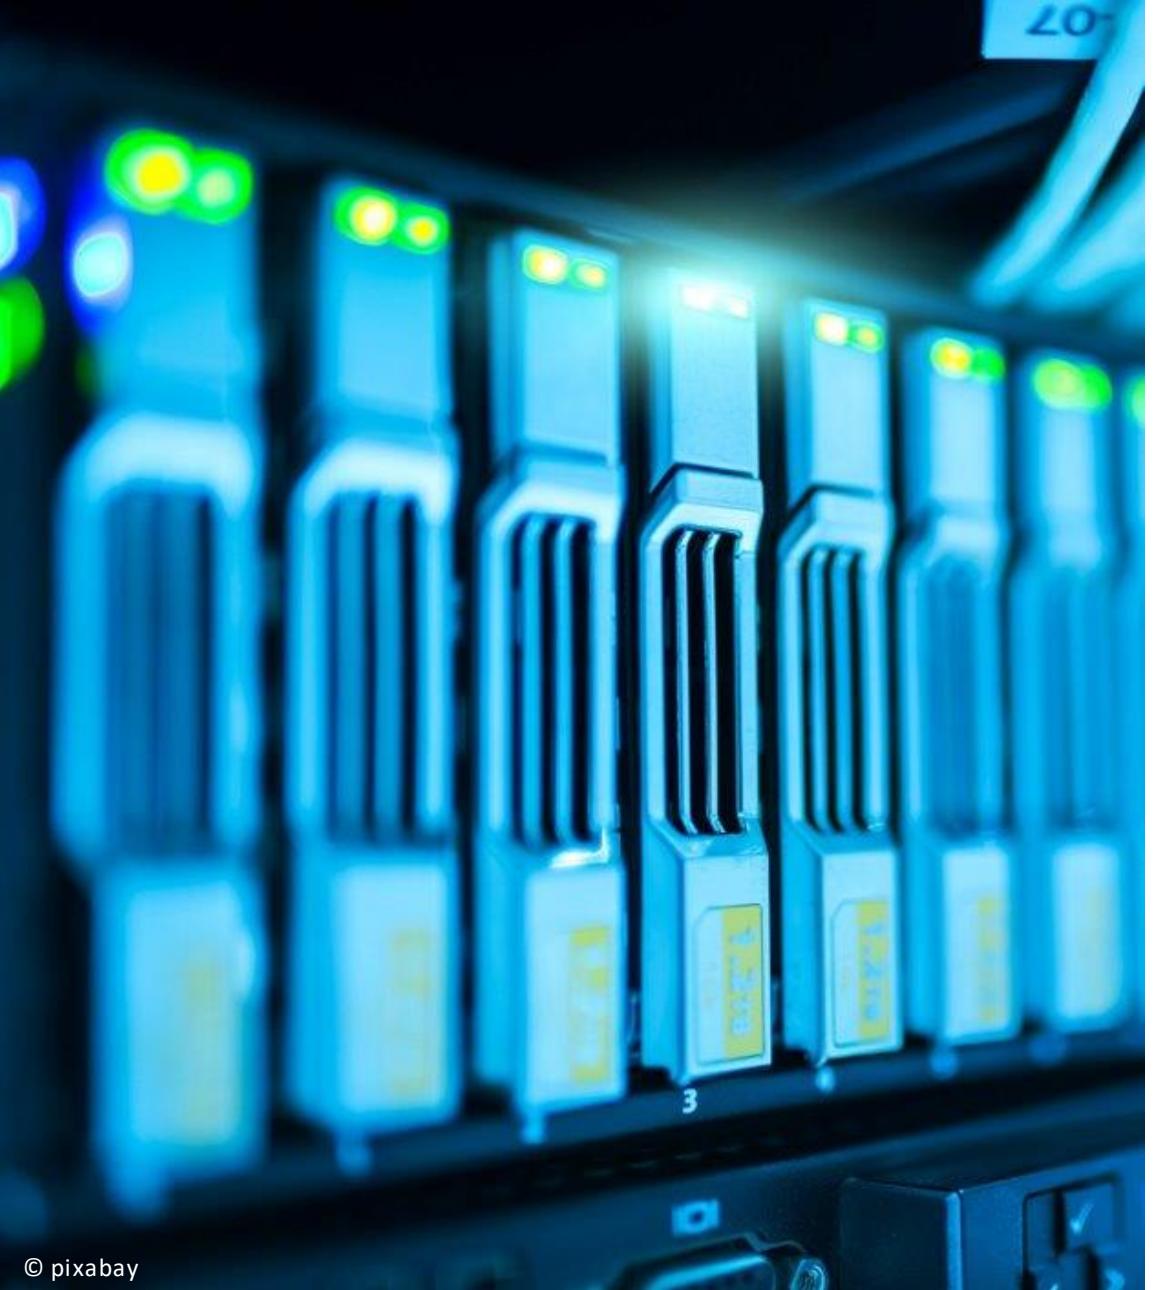
BeeGFS<sup>[3]</sup>

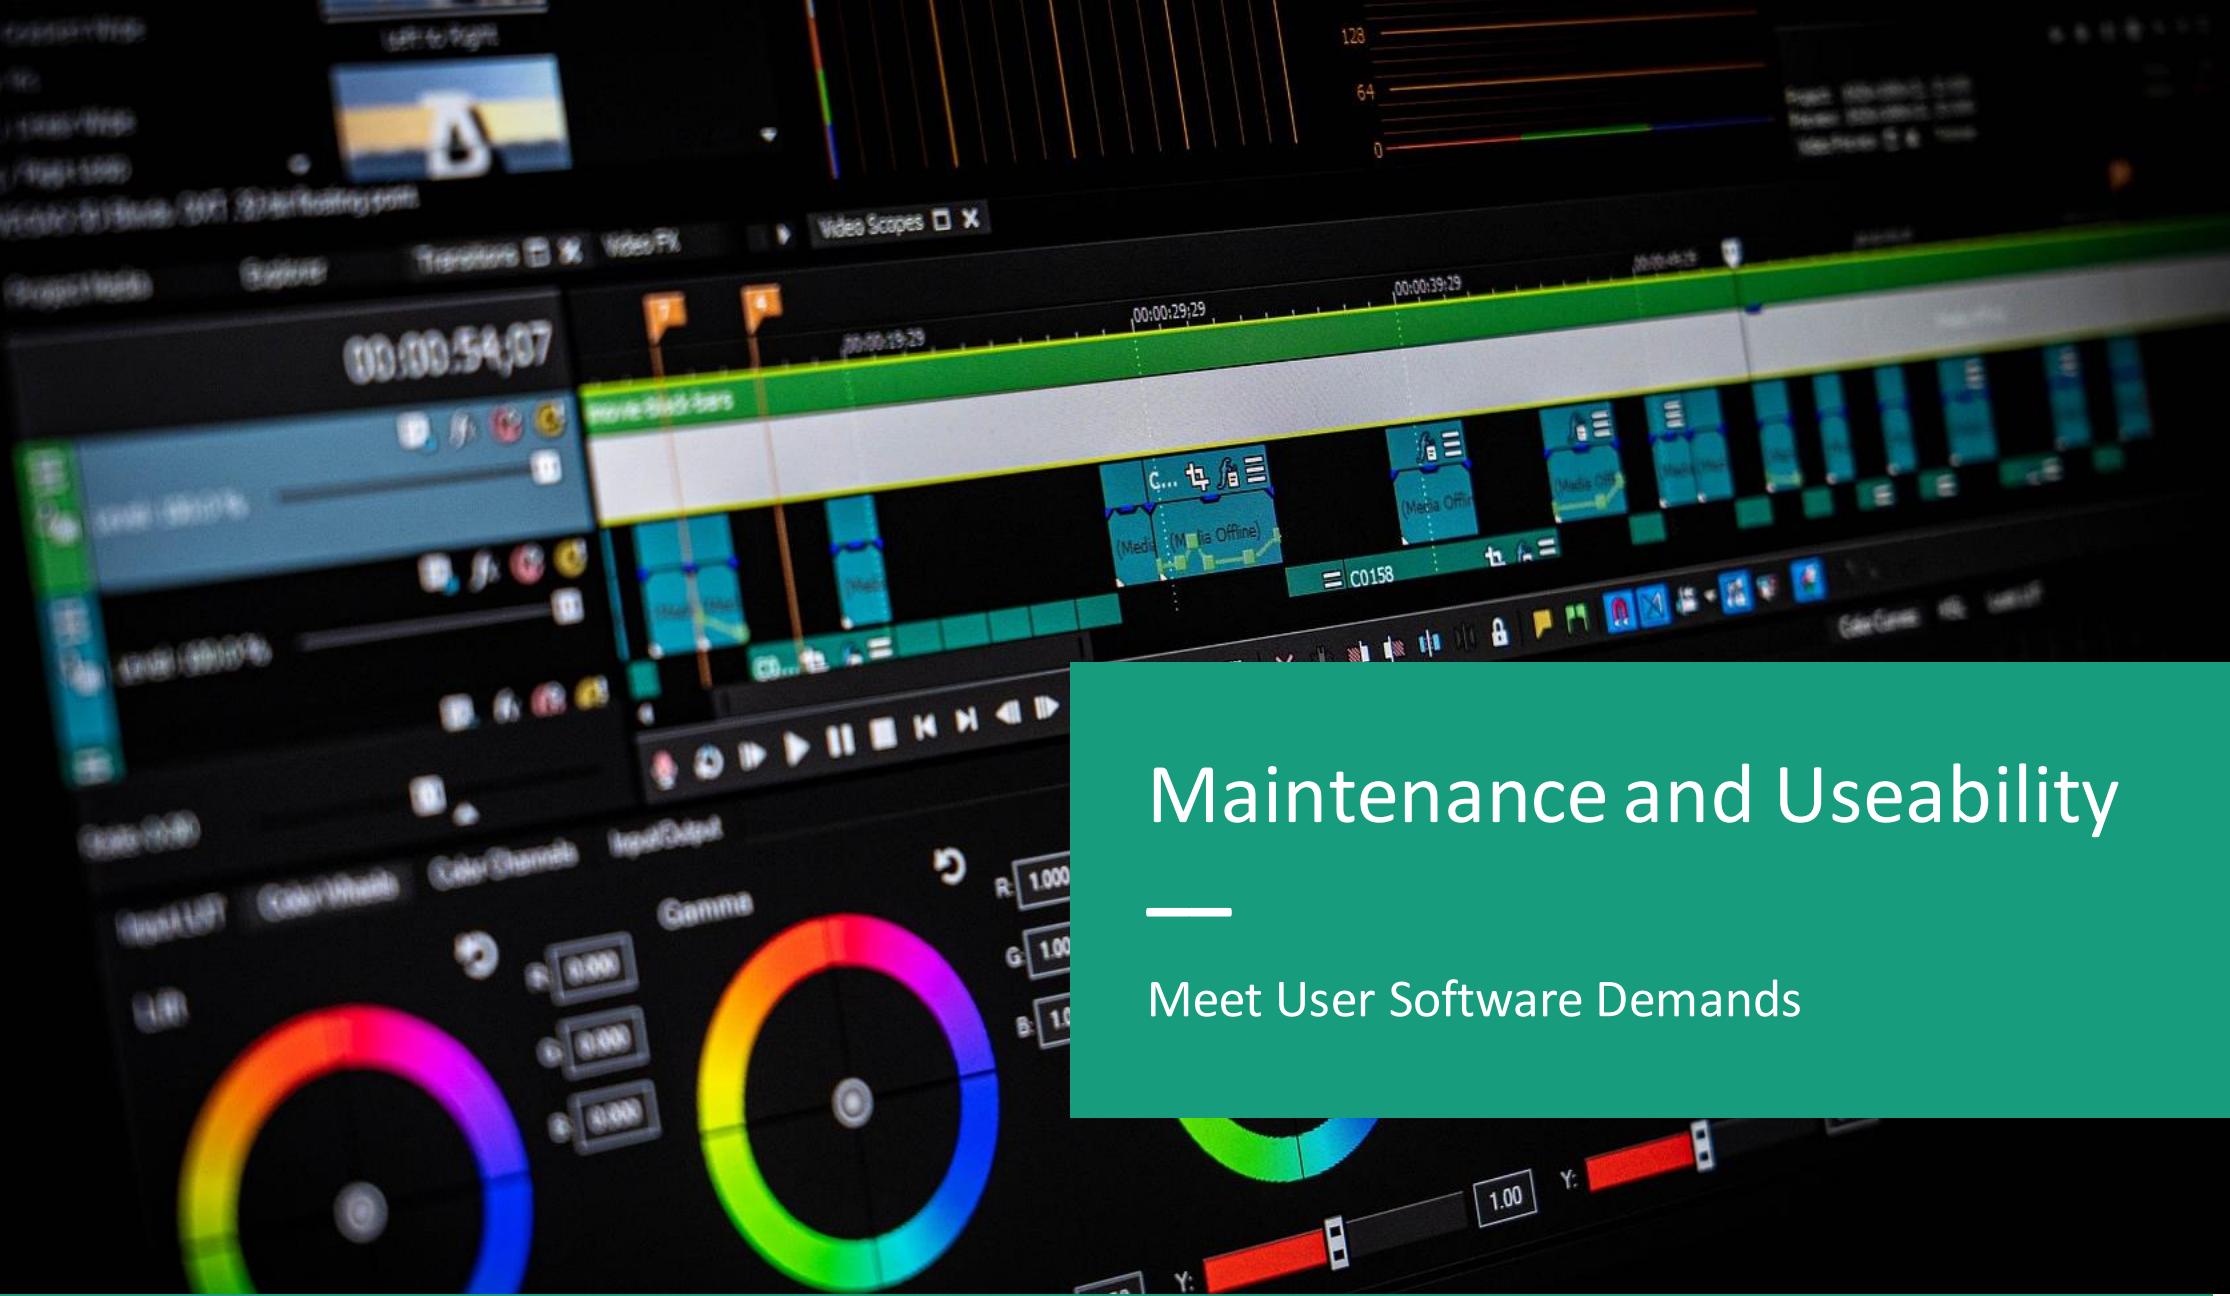


- redundant, parallel file system
- **huge data throughput** (via infiniband)
- temporary job FS (BeeOND) on local SSDs between node
- **open-source** (commercial support available)

[3] <https://www.beegfs.io>

© pixabay





## Maintenance and Usability

—  
Meet User Software Demands

# Maintenance and Usability

## Demandings we have to care

### ■ using graphic cards for calculations

- install drivers, the corresponding libraries and TensorFlow, PyTorch, ...
- easy GPU access

### ■ ML/DL algorithms

- libraries, programs and dependencies change fast
- different algorithms need different programs/libraries

### ■ programs/tools often available as deb-packages

- but many HPC clusters are rpm-based

### ■ HPC clusters have very heterogeneous users

- need a lot of different tools and some dependencies may collide
- users cannot be root



© pixabay

# Maintenance and Usability

## Easy for Users and Admins

- **easy to use** (for users)
  - access via **web interface**
  - **interactive development tools**  
(e.g. TheiaIDE, JupyterLab)
  - direct **GPU access**
  - uncomplicated **multi-node/-GPU usage**
  
- **easy maintainable** (for administrators)
  - provide different libs for specific user groups
  - use **anaconda environments**
  - stick to **singularity containers**



© pixabay / Fraunhofer ITWM

# Maintenance and Usability

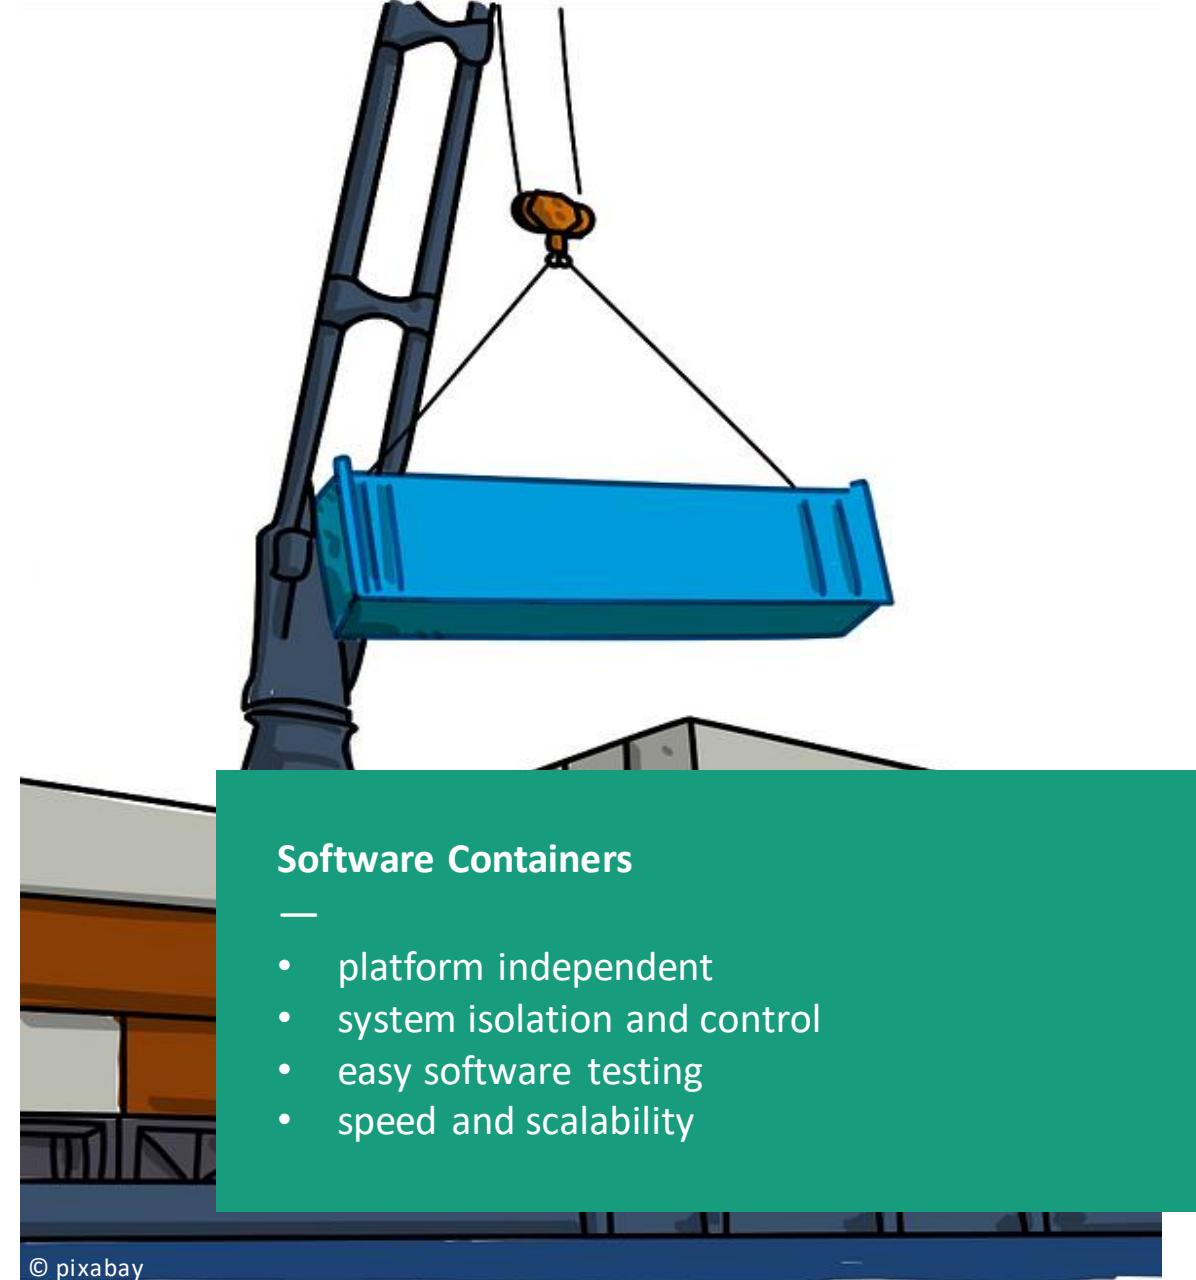
## Portable Software Environments

### singularity<sup>[4]</sup> containers



- provide the **OS** that ML/DL users **need**
- **on the host:** base OS and graphics drivers  
**in the image:** all other dependencies
- **no root-privileges** to start the container
- **only needed folders** mounted  
(e.g. /home/USERNAME, /scratch/USERNAME)
- create **images from scratch** (takes a few minutes)  
or **transform** an **existing** (clean) **docker** image

[4] <https://www.sylabs.io/singularity>



© pixabay

# Maintenance and Usability

The base Environment Carme provides

## Carme Singularity "Base Image"

- set of **basig software**  
(e.g. vim, nano, gcc, gdb, ... )
- everything **ready to use IB**  
(using free drivers)
- **MPI** and **GPI** are installed per default  
(GPISpace installation ready)
- **miniconda** and/or **mambaforge**
- **Python 3.X**
- **TheiaIDE** and **JupyterLab**
- **headless or web-based applications** can run inside the container

# Maintenance and Usability

## Userspace Package Management

### anaconda<sup>[5]</sup> based environments

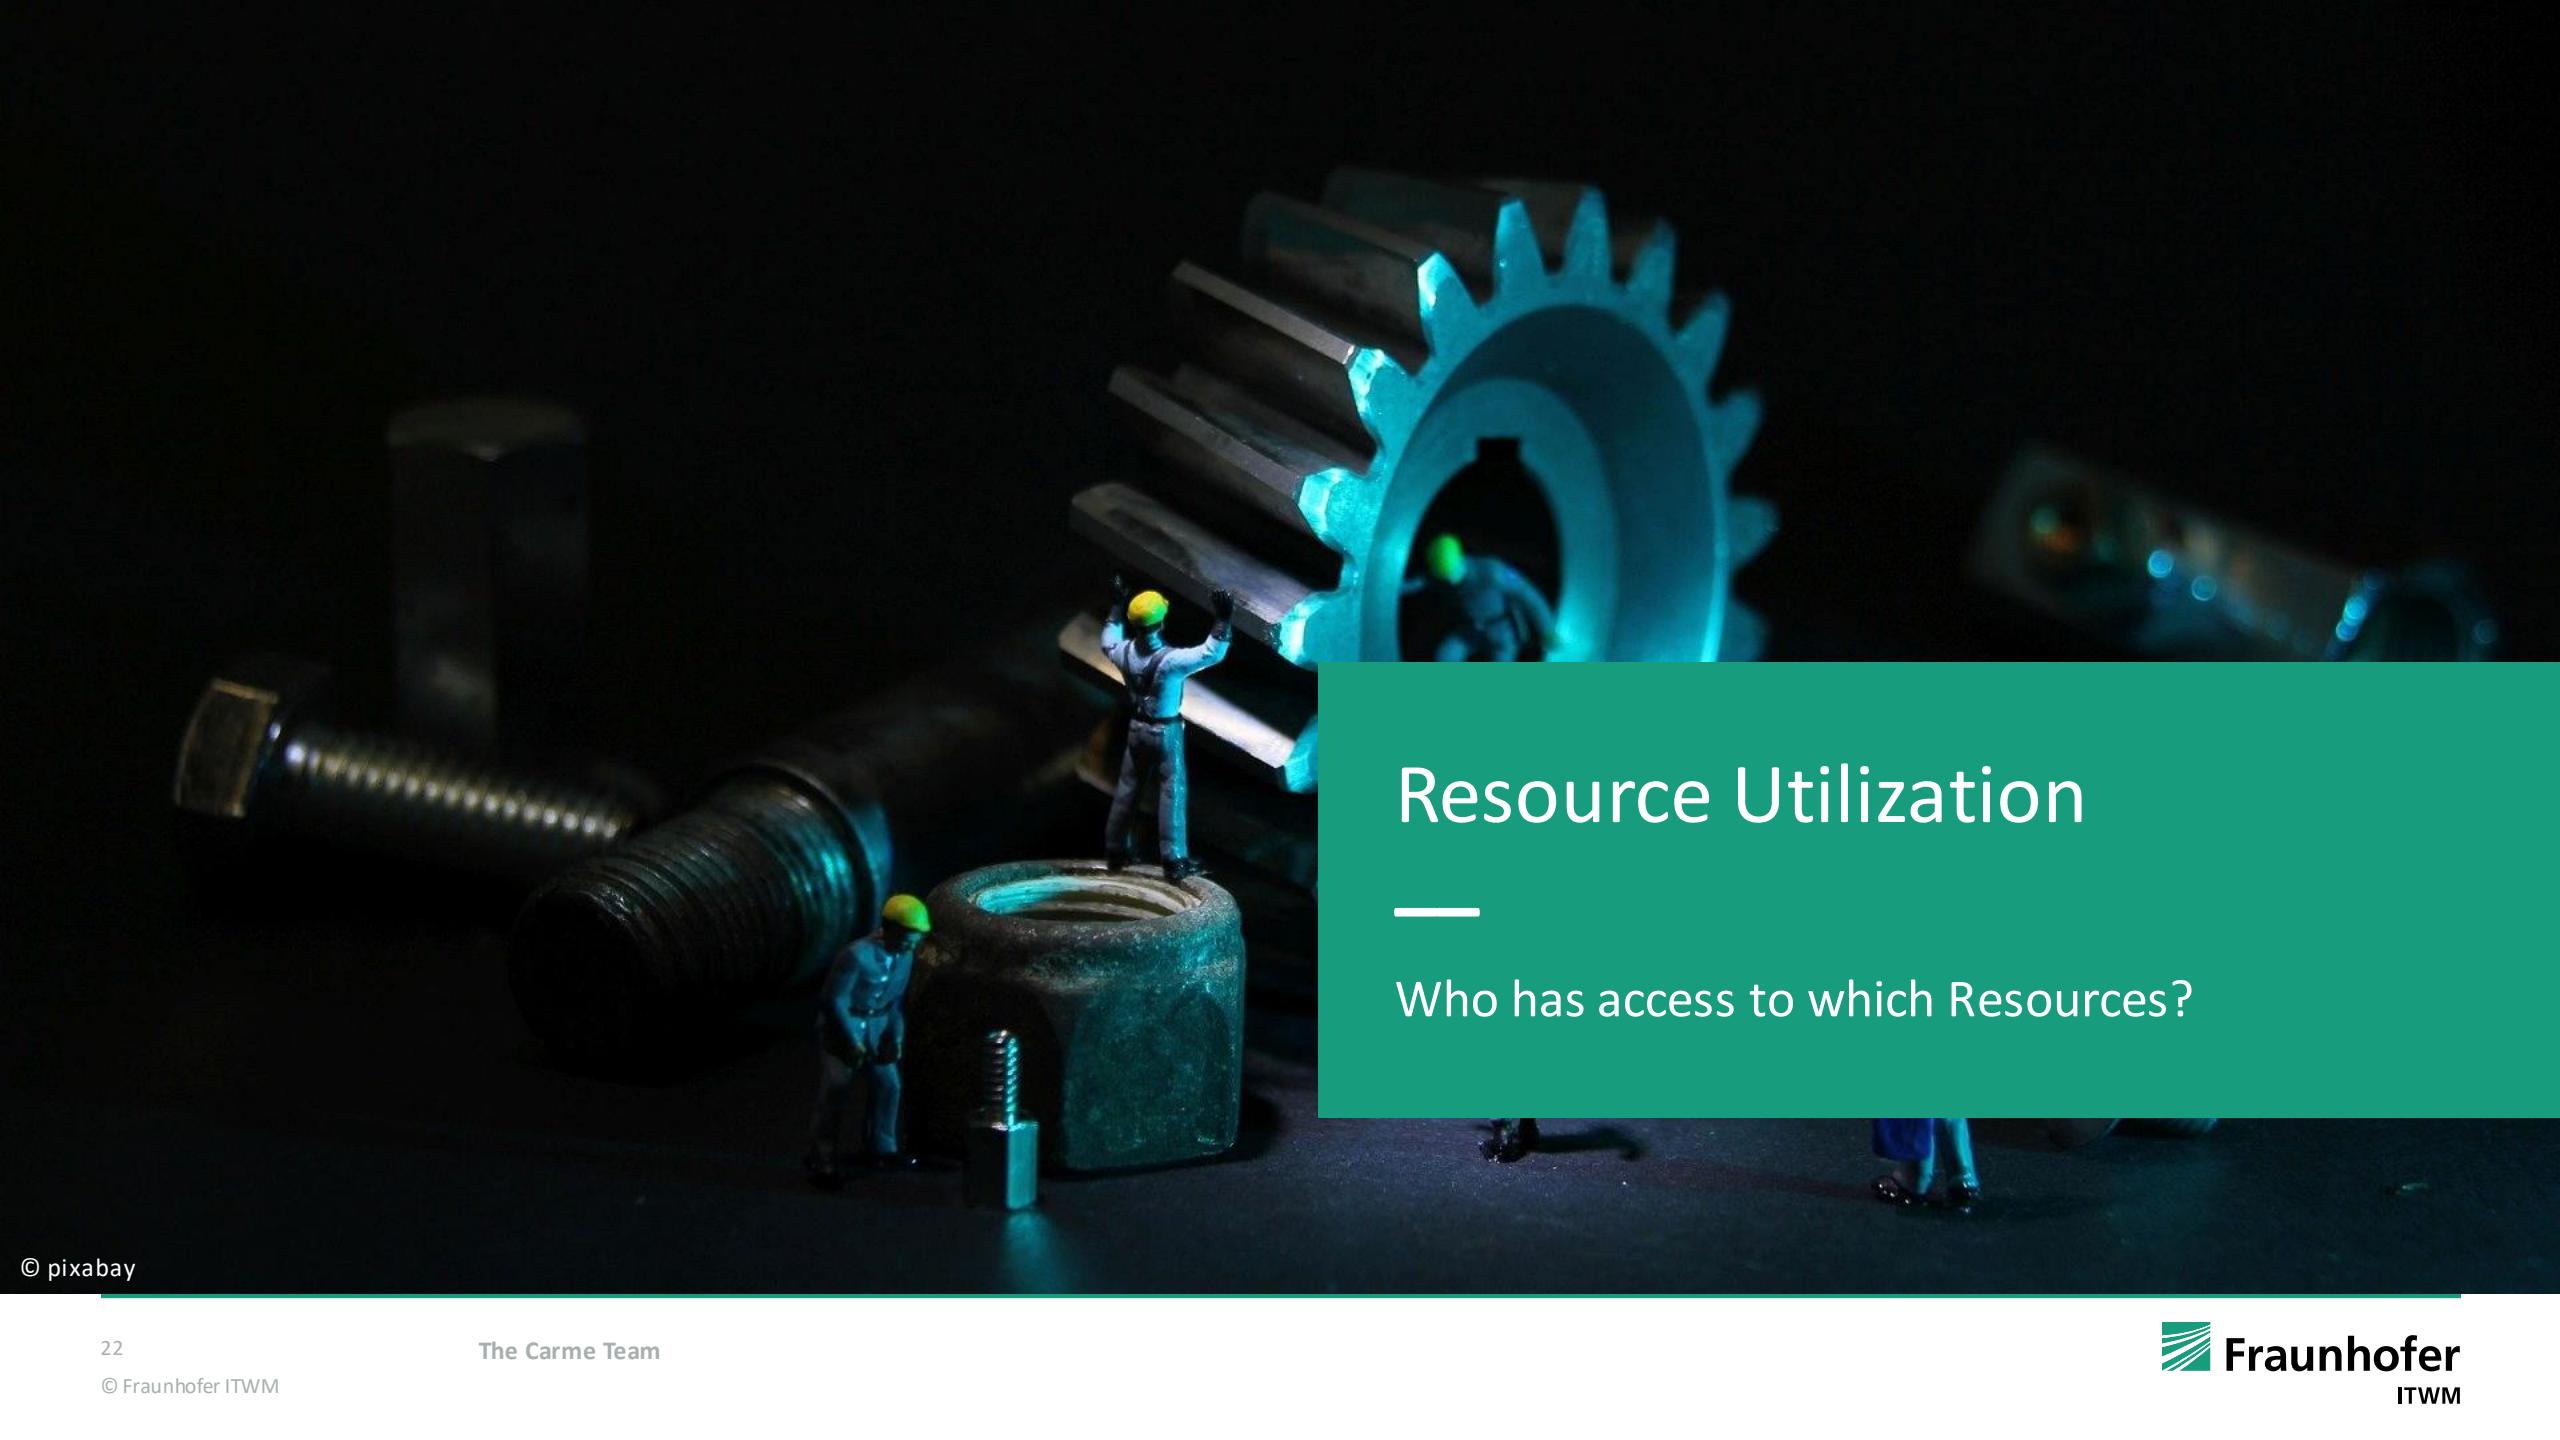


- **open-source** (normal and enterprise versions)
- central **miniconda/mambaforge** installation inside the image (people can use it directly)
- users have the freedom to **install** (most) of the **libs and tools** they need (precise version control)
- all **PyTorch** and **TensorFlow** versions possible
- easily **share** and **save** your environments
- supporting **cooperations**

[5] <https://www.anaconda.com>

© pixabay





## Resource Utilization

—

Who has access to which Resources?

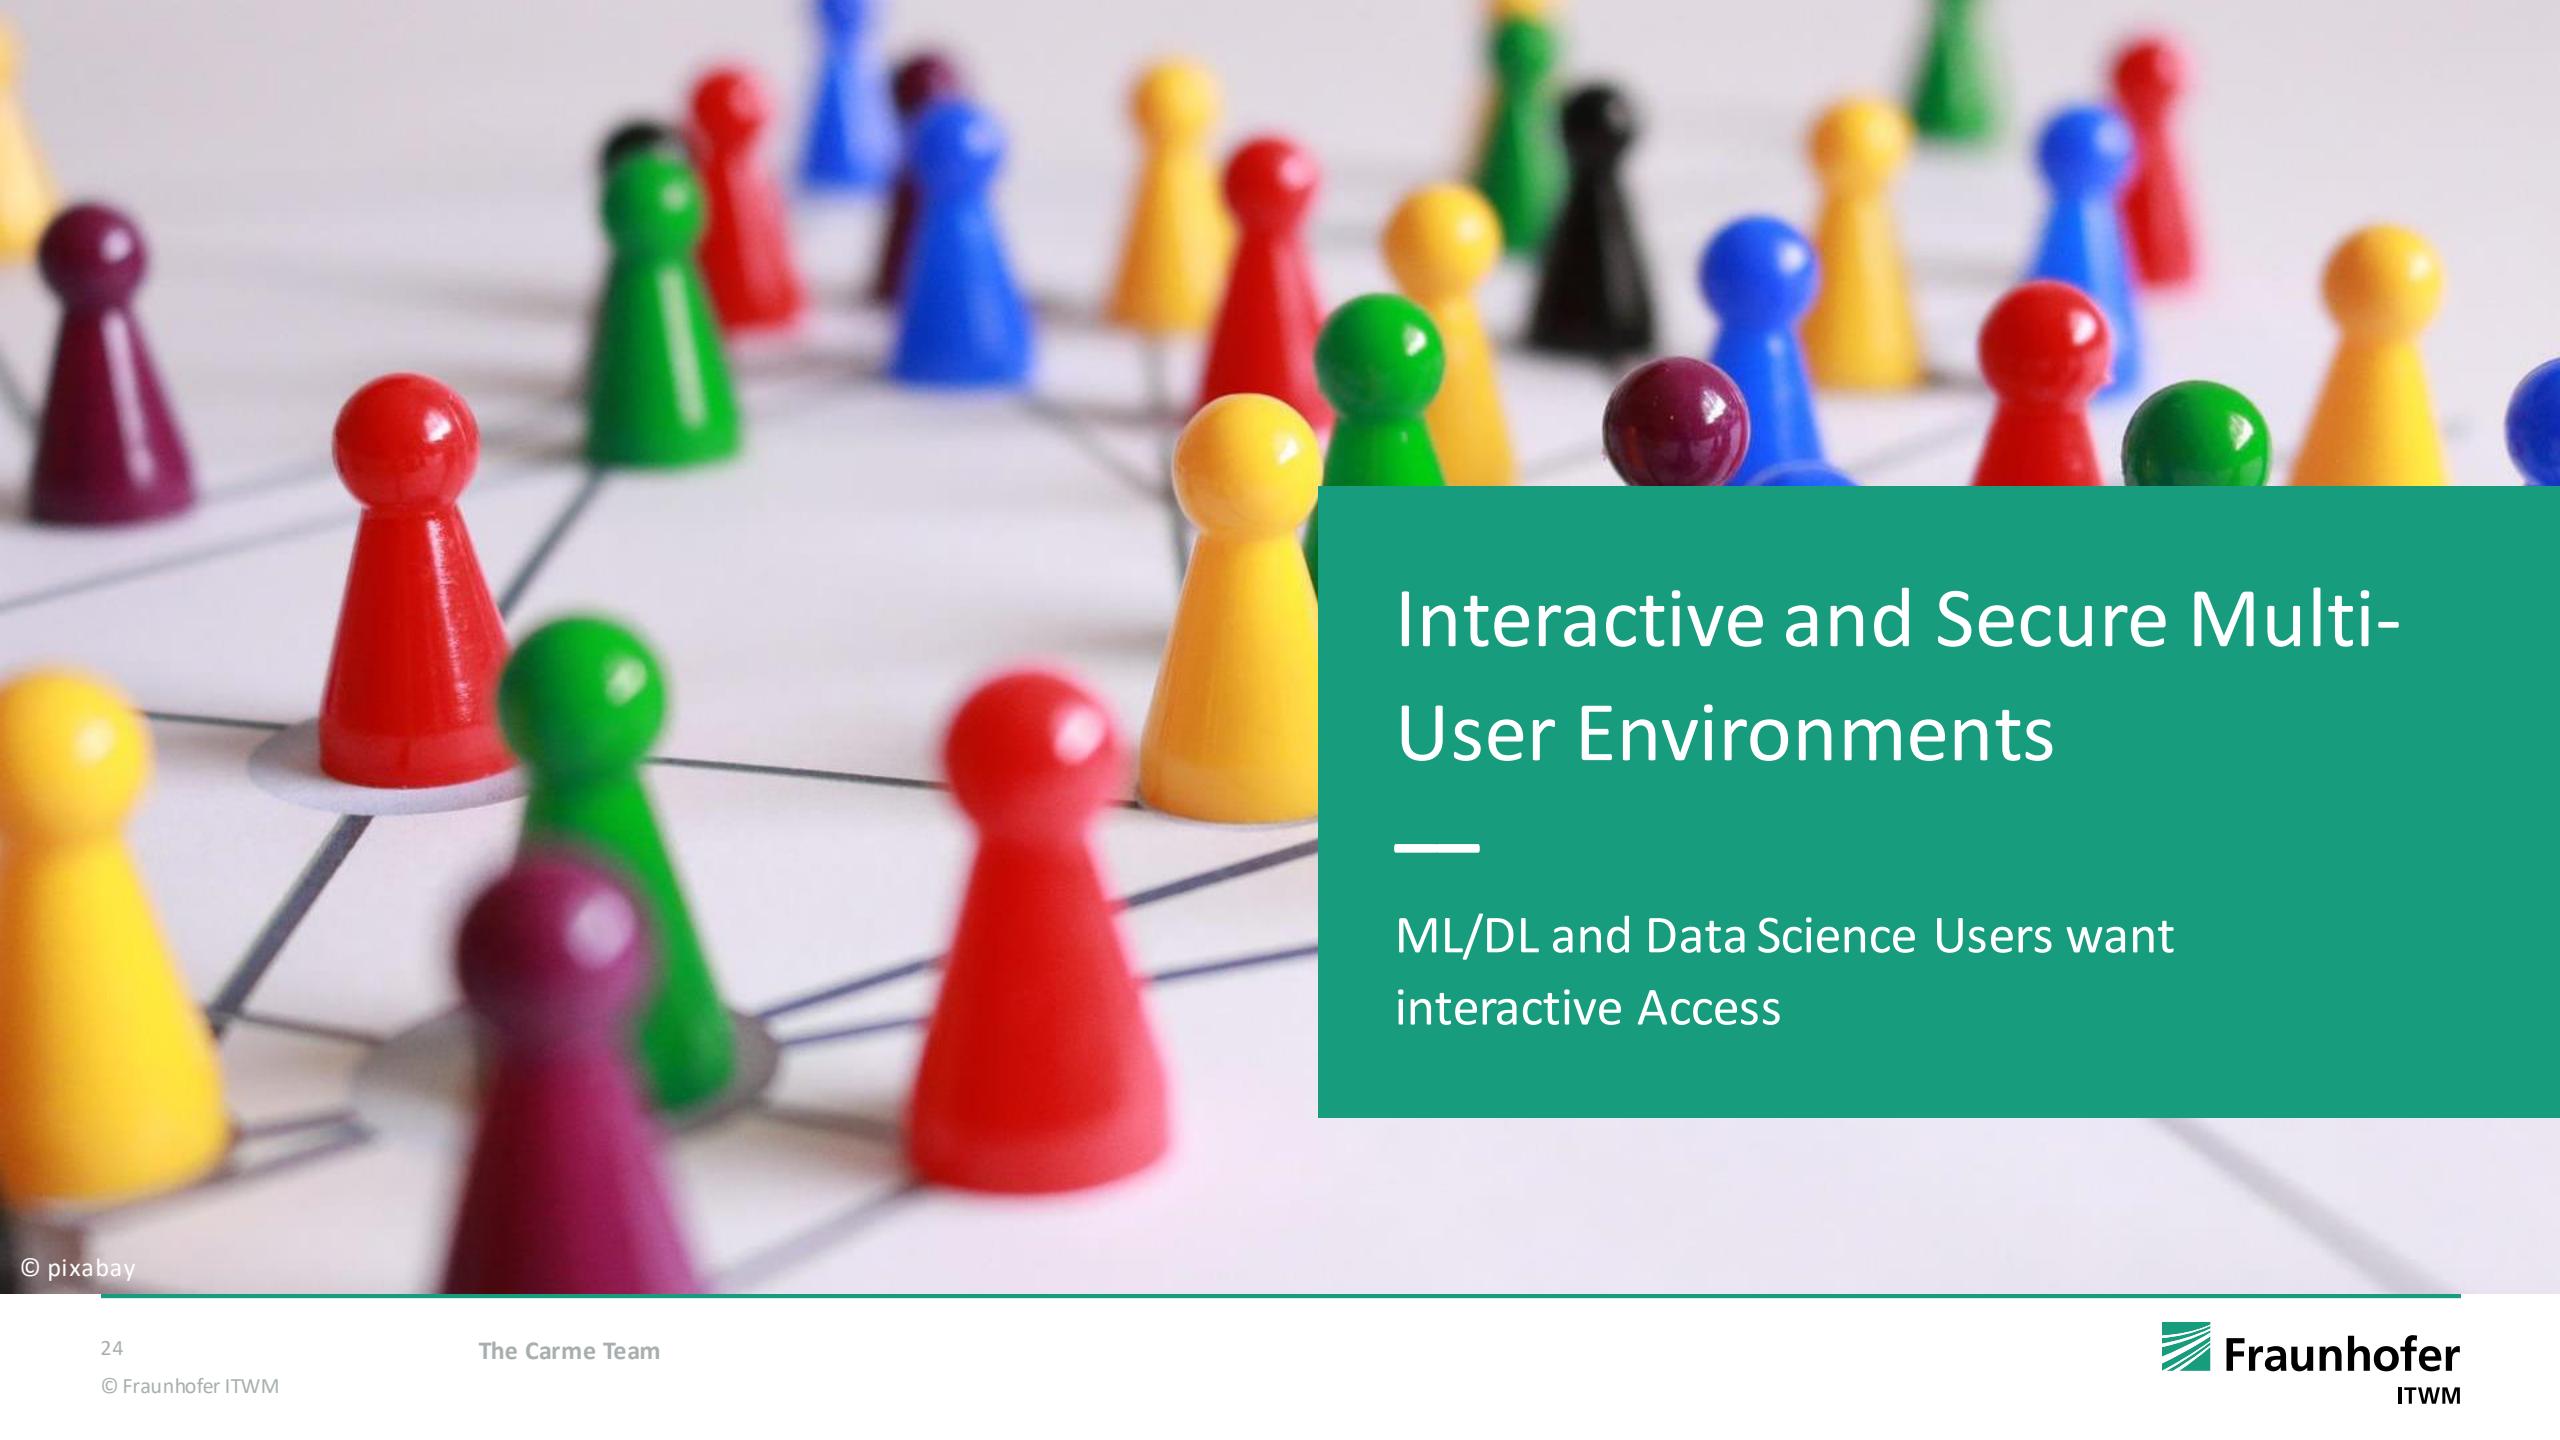
# Resource Utilization

Control the Resources you have

- CARME utilizes the **scheduler resource limits**  
(fully respects non-carme jobs, cgroups, queues, etc.)
- **additional limits** can be defined **on top** of the scheduler limits  
(modify #(nodes), #(GPUs), running jobs, etc.)
- **integrates smoothly in your existing scheduler**  
(only a carme feature flag is needed)
- everything **inside a job** runs inside an **isolated namespace** (programs cannot escape a job)
- inside a job you can only see the GPUs associated with the job (no access to information of other GPUs on the node)



© iStock



# Interactive and Secure Multi-User Environments

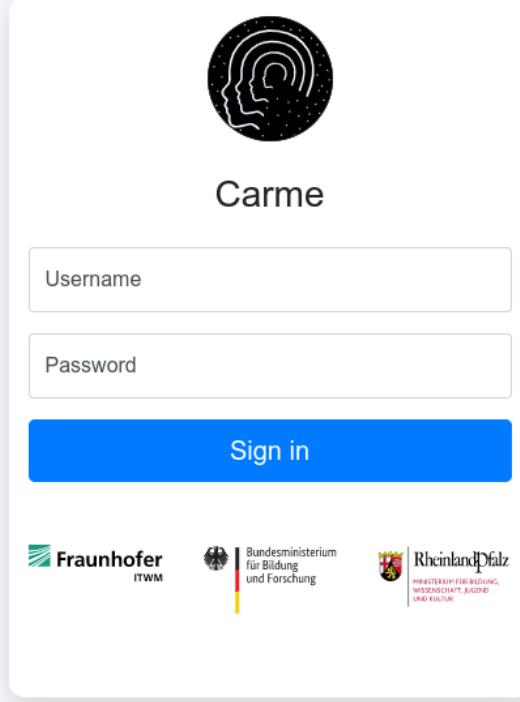
---

ML/DL and Data Science Users want interactive Access

# Interactive Usage

## Secure and simple Access

- access via **web interface**
- **no additional software**  
(like ssh or special tool)
- runs **behind your firewall**
- **secure authentication**  
(e.g. via LDAP) and additional  
**2FA**



The screenshot shows the Carme web interface login screen. At the top center is a circular logo featuring a stylized profile of a head with concentric circles around it. Below the logo, the word "Carme" is written in a sans-serif font. Underneath the name are two input fields: one for "Username" and one for "Password", both represented by light gray rectangular boxes. Below these fields is a large blue rectangular button with the white text "Sign in". At the bottom of the page, there are three logos: "Fraunhofer ITWM" (with a green square icon), "Bundesministerium für Bildung und Forschung" (with the German eagle logo), and "Rheinland-Pfalz MINISTERIUM FÜR BILDUNG, WISSENSCHAFT, JUGEND UND KULTUR" (with the coat of arms of Rhineland-Palatinate). At the very bottom of the slide, there is a thin green horizontal bar containing the text "Disclaimer" and "Data Privacy" on the left, and social media icons for Twitter, LinkedIn, and YouTube on the right.

# Interactive Usage

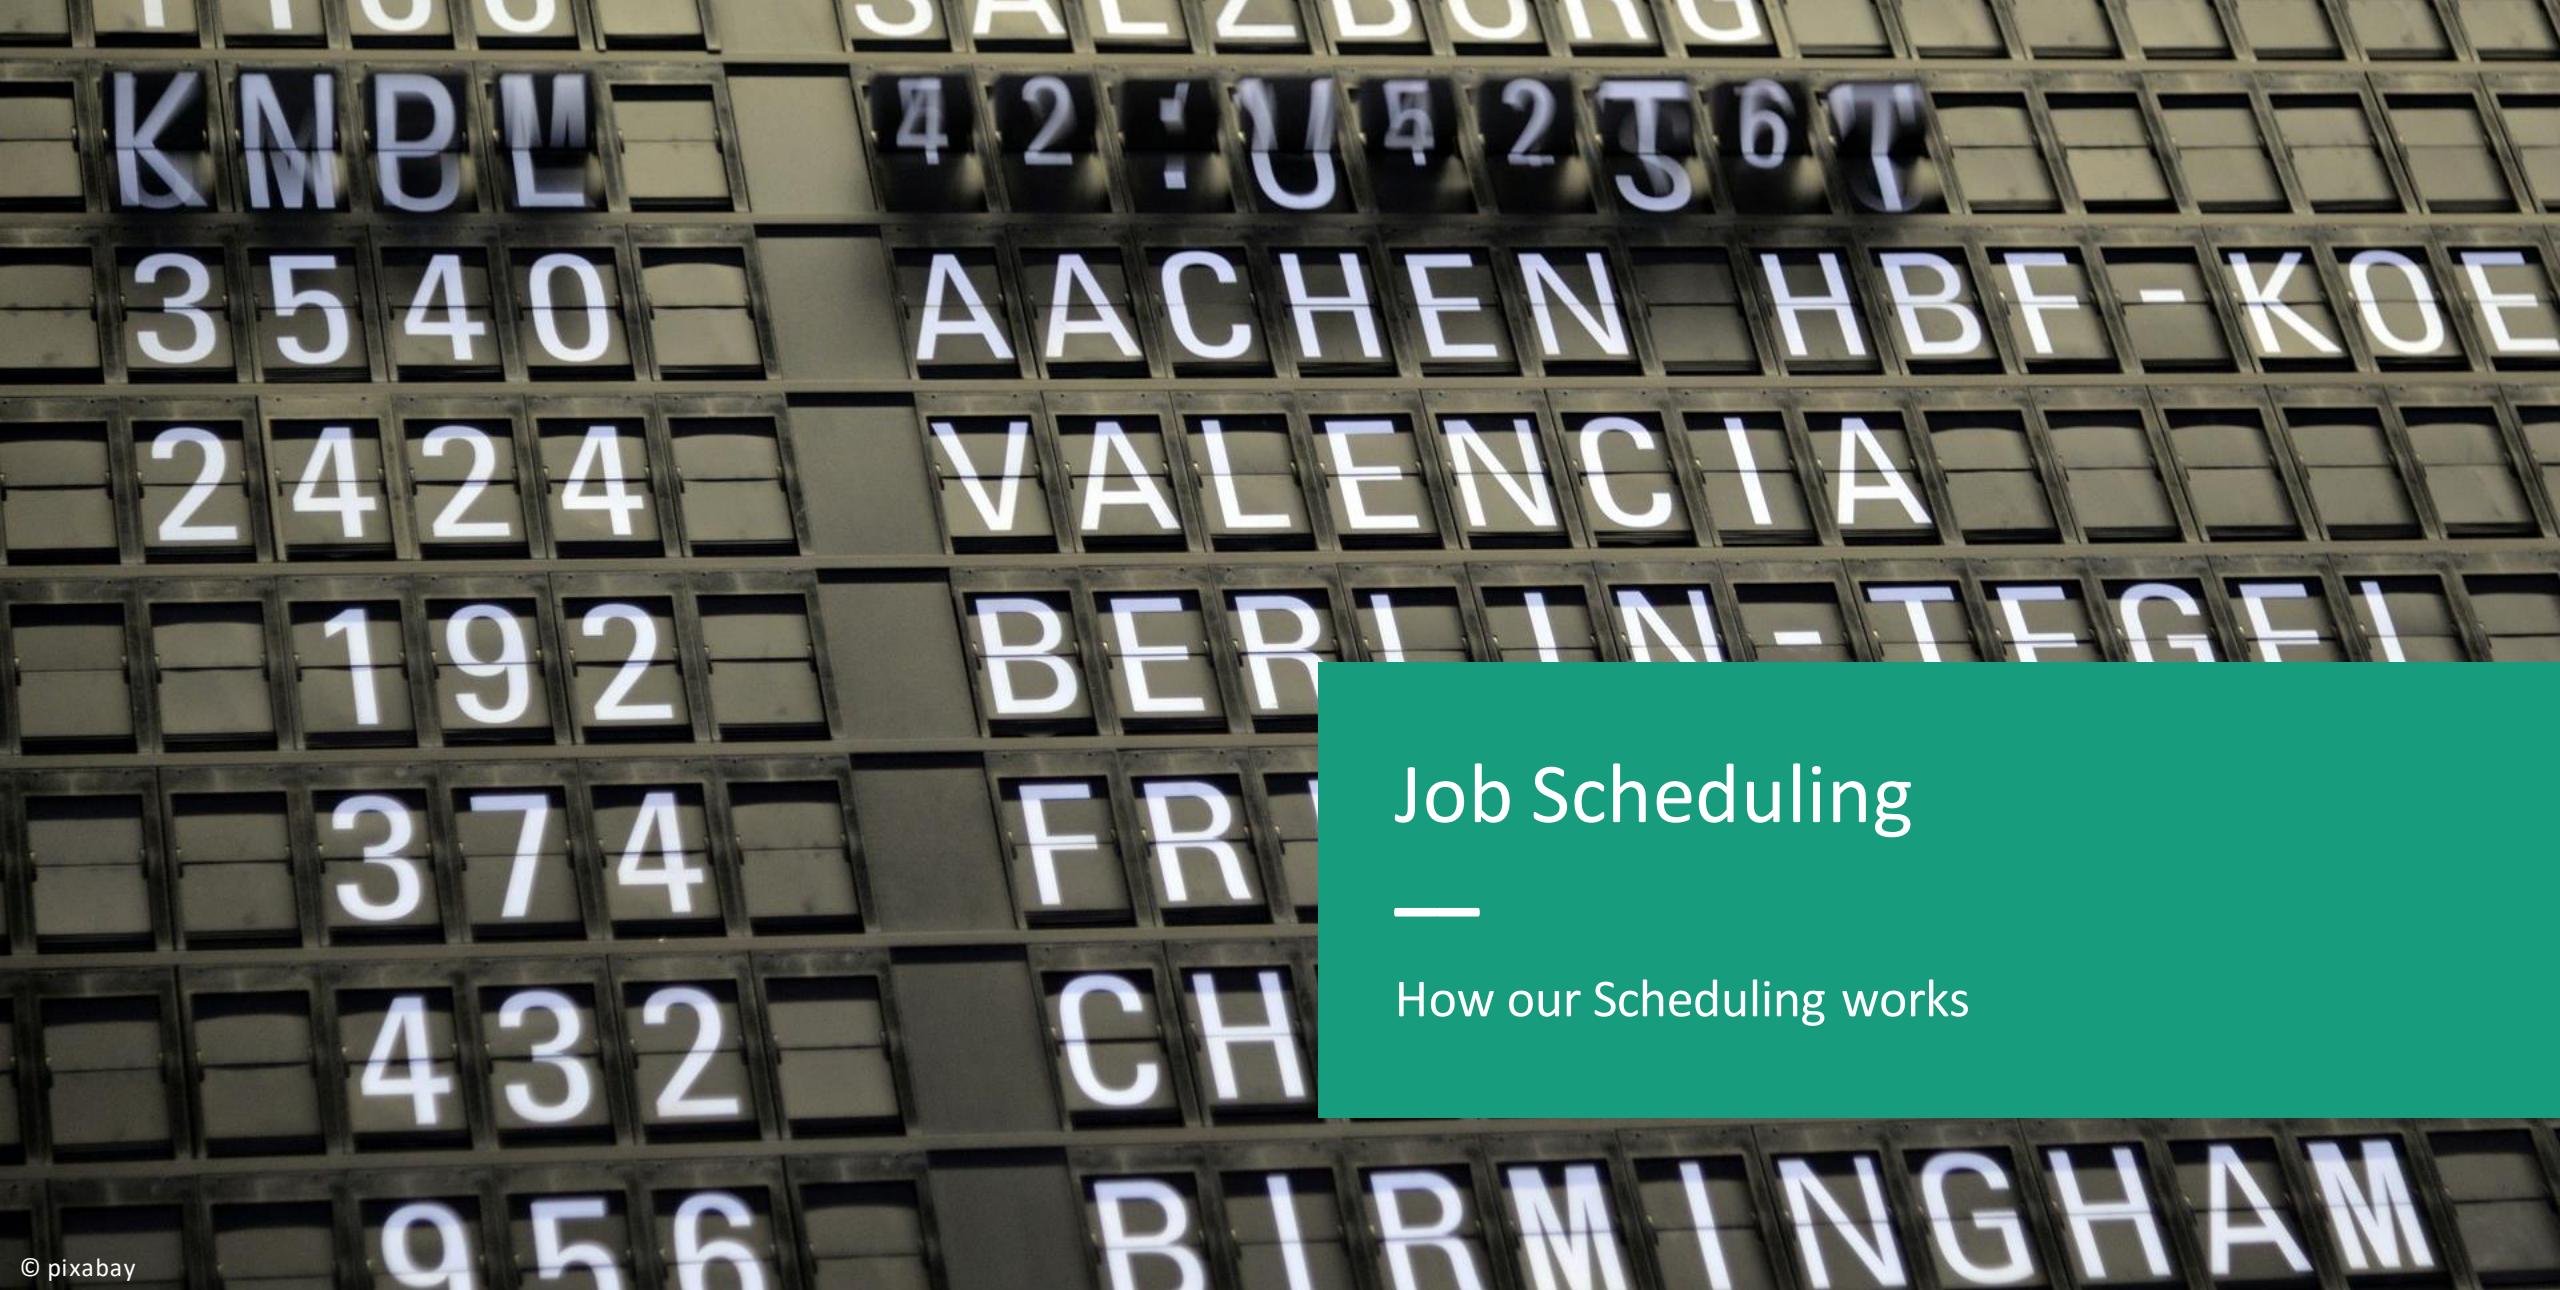
## The User Control Center

two clicks to development

(1) start a job

(2) select your preferred development environment

The screenshot shows the Carme User Control Center dashboard. At the top, there's a navigation bar with links for Dashboard, History, Help, and user demo-admin. Below the navigation is a news section titled "Welcome to the STYX". To the right of the news is a GPU status card showing "Total GPUs: 108" with 32 GTX GPUs and 76 OTHERS. It also displays metrics for GTX: 2 GPUs/node, 10 CPUs/GPU, and 30 GB RAM/GPU. To the right of the GPU card is a chart titled "Chart" showing GPU usage over time (CET). The chart has tabs for Total, GTX, TITAN, and A100. It tracks Free (green circles), Used (blue squares), and Queued (black diamonds) GPUs. The chart shows a slight upward trend in used GPUs over the period from Now to April 26, 16:29. Below the chart is a "Jobs" section with a table of running jobs. The table includes columns for ID, Name, Image, Nodes, CPUs/node, GPUs/node, Status, Timing, Entry points, Details, and Stop. Job 34998 is for CarLA with 1 node, 16 CPUs, 2 GPUs (TITAN), and is marked as "✓" (running). Job 34999 is for AlphaGo with 4 nodes, 96 CPUs, 4 GPUs (A100), and is marked with a gear icon (pending). The dashboard also includes links for Disclaimer, Data Privacy, and footer text "© 2018-2022 Carme".

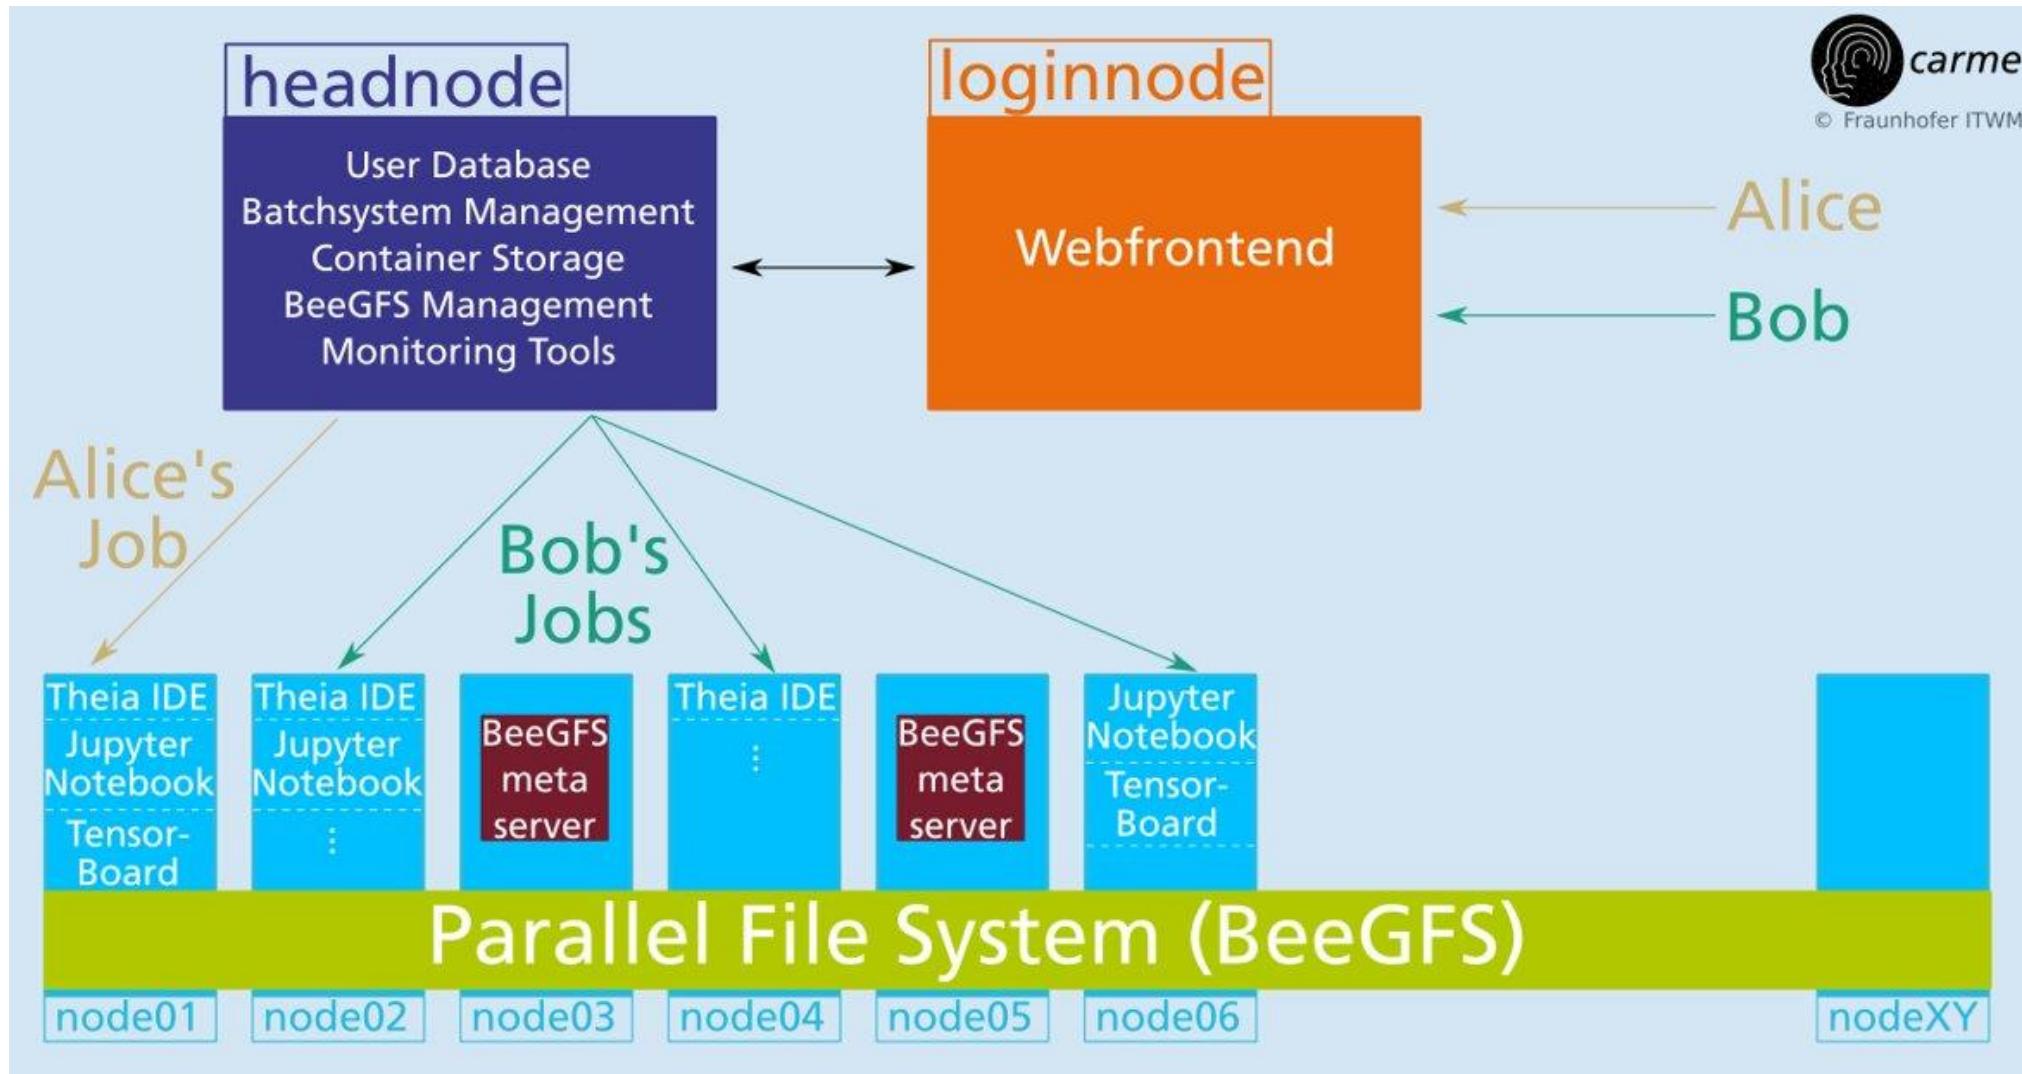


## Job Scheduling

---

How our Scheduling works

# Job Submission Scheme



# Job Submission Scheme

- user login to the web interface
- user submits a job
  - request for new job is sent from the frontend to our hpc-backend
  - secure connection via rpyc and special keys
- our hpc-backend
  - gets the job information from the frontend and our data base
  - finally submits the job to the scheduler as original user
- job start script is handed over to the scheduler
  - scheduler checks handed over information  
(resource limits, user accounts, availability, ...)
  - respective CARME prolog scripts run  
(only if the job belongs to us)
- job starts on the requested node(s)
- user is informed in the frontend that the job has sucessfully started



© pixabay

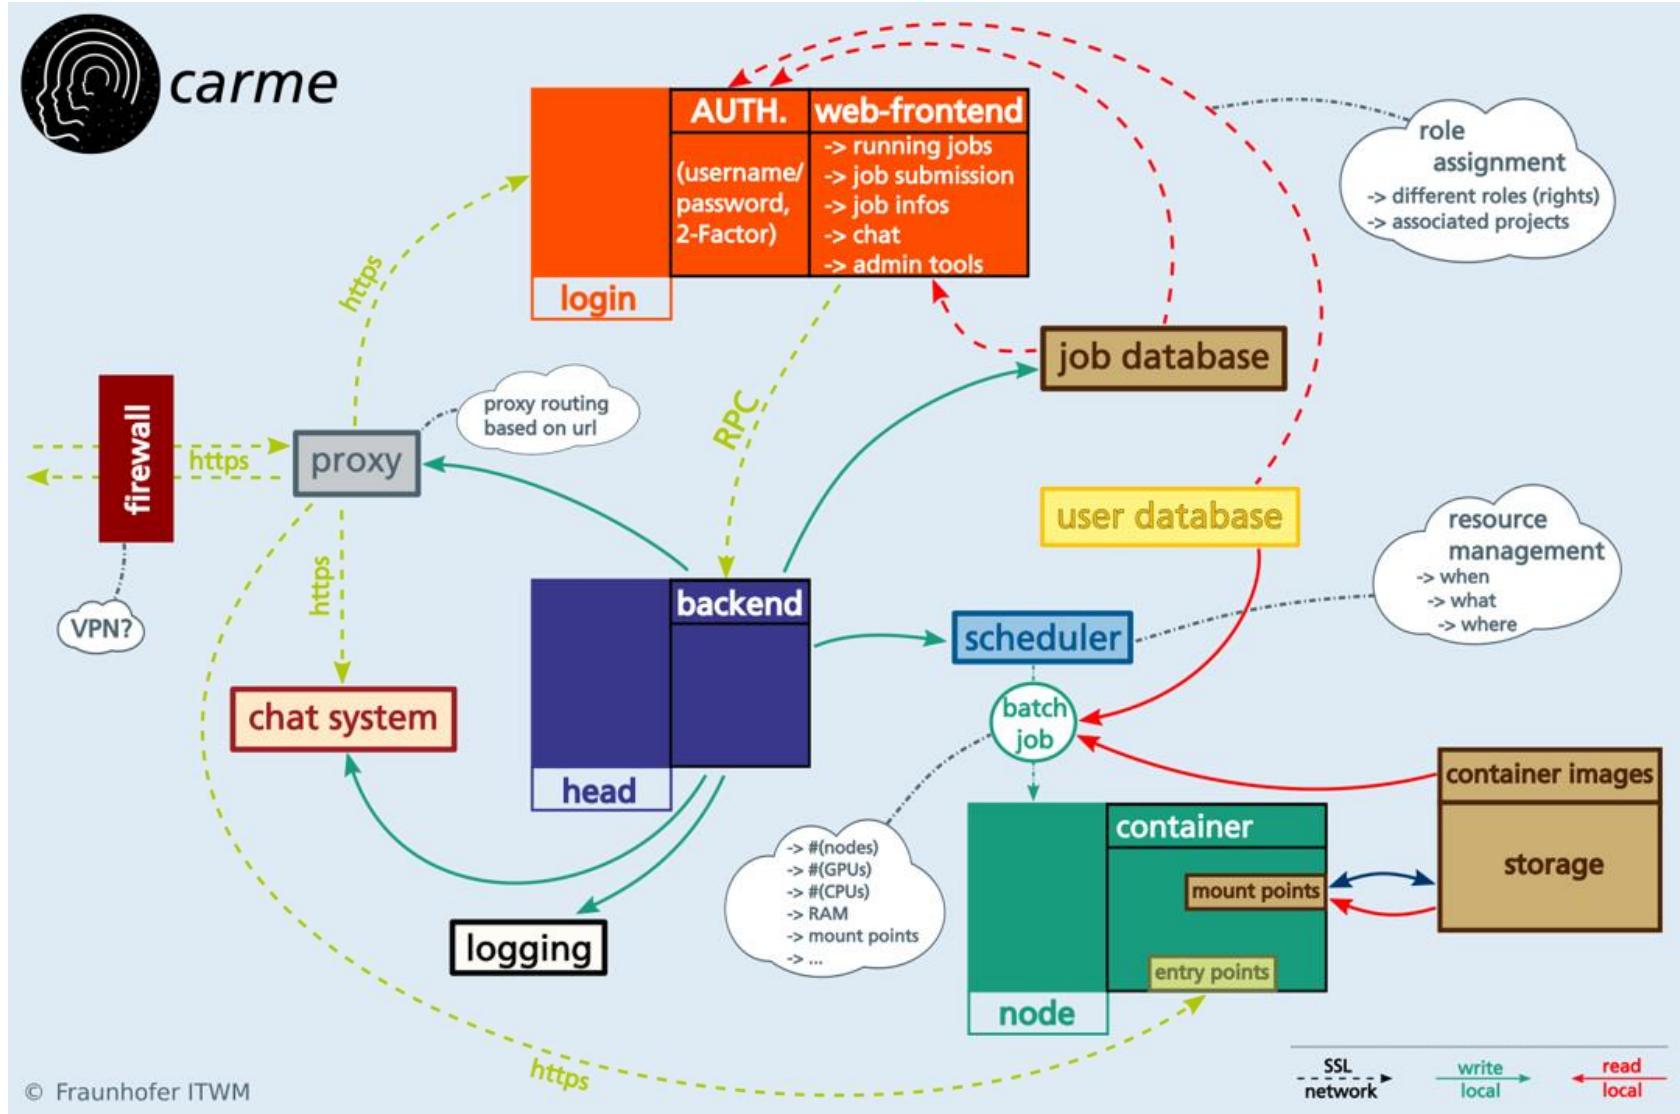


# Security Concept

---

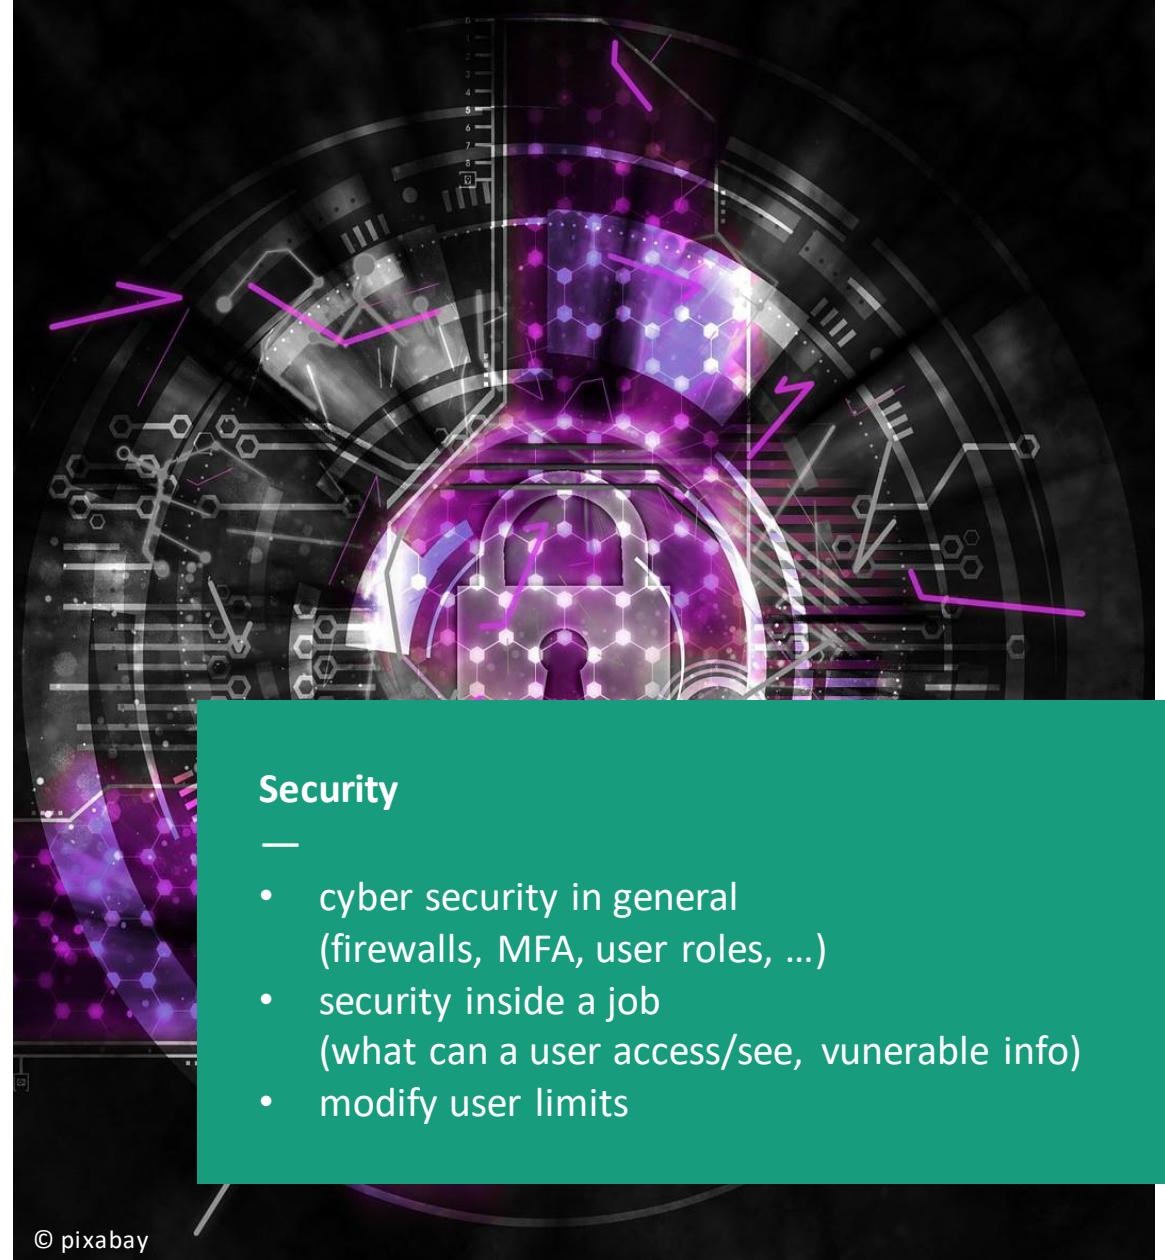
Our Security Idea

# Security Concept



# Security Concept

- everything can run **behind your firewall**
- we utilize our **own proxy and frontend**  
*(both are separated singularity containers)*
- **secure connection** between frontend and backend
- access **requires authentication** and **2FA**
- define what is **mounted inside a job**
- fully **integrate the limitations** of your **scheduler**
- **additional limitations** in CARME  
(e.g. resource limits, running jobs and singularity images)
- everything **inside a job** runs in its **own namespace**  
(no chance to use other resources)



© pixabay

## Security

- 
- cyber security in general  
(firewalls, MFA, user roles, ...)
- security inside a job  
(what can a user access/see, vulnerable info)
- modify user limits



## Our Key Features

---

In a Nutshell

# Our Key Features

## In a Nutshell

- **seamless integration** into existing HPC setups
- user-friendly **web-interface** providing **flexible** and os-**independent** access from anywhere in the world
- **interactive jobs** to develop directly **on the cluster** with your favorite **deep learning tools**
- fully **separated jobs** with **custom resources**
- **intuitive abstraction** of complex cluster topologies
- **distributed multi-node/multi-gpu jobs** with direct access to GPI, GPI-Space, MPI, Tarantella and **Horovod**
- user maintained and **containerized environments** using singularity and **anaconda**

# Our Key Features

## In more Detail

### ■ **seamless integration with available HPC tools**

- job scheduling via SLURM
- native LDAP support for user authentication
- integrate existing distributed file systems like BeeGFS

### ■ **access via web-interface**

- OS independent (only web browser needed)
- full user information (running jobs, cluster usage, news, messages)
- start / stop jobs within the web-interface

### ■ **interactive jobs**

- flexible access to GPUs
- access via web driven GUIs like Theia-IDE and/or JupyterLab

### ■ **distributed multi-node and/or multi-gpu jobs**

- easy and intuitive job scheduling
- directly use GPI, GPI-Space, MPI, Tarantella and Horovod within the jobs

### ■ **full control about accounting and resource management**

- job scheduling according to user specific roles
- compute resources are user exclusive

### ■ **user maintained, containerized environments**

- singularity containers  
(runs as normal user, GPU, Ethernet and Infiband support)
- anaconda environments  
(easy updates, project/user specific environments)
- Built-in matching between GPU driver and ML/DL tools



## In Summary

---

What you get

# In Summary

## What you get

### The User Side

- **interactive development** access to HPC systems  
(keep your workflow with Jupyter Notebooks or full IDE)
- **easy way to submit and get jobs** on big clusters  
(simplified submission mask, no need to write batch scripts)
- guarantee that **only you** have **access to your resources**  
(no need to worry about other users)
- utilize **containers** and **conda environments** to get software  
(keep track with quickly chaning and diverse tool requirements)
- **access your HPC cluster** via a **web browser**  
(no need to install special or additional software)

# In Summary

## What you get

### The Administrator Side

- ▶ **easily integrate** CARME in your scheduler set up  
(a feature flag and our prolog/epilog scripts is all you need)
- ▶ **control** which **folders** user have access to  
(you can decide per image which folders are mounted or not)
- ▶ **provide the basic software** that ML/DL users need to do there work  
(keep the host installation small and clean)
- ▶ manage our **data base** via a simple **web interface**
- ▶ **simplify the support** for AI and Data Science users
- ▶ get **logs** and **error outputs** for CARME scripts and services

# In Summary

## What you get

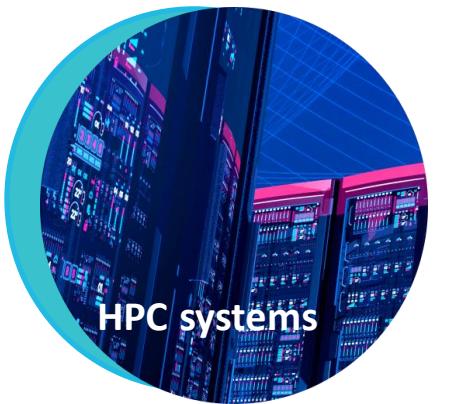
### The Institution Side

-  CARME was and will be **developed in Germany**  
(following Fraunhofer standards and respecting data privacy)
-  license that allows **commercial use**
-  giving users a tool to **easily utilize GPU resources**  
(be sure that your investments in hardware pay off)
-  **increase the productivity** of both your AI developers  
and researchers and your system administrators

# Bringing AI Development / Education and HPC Systems together

## The CARME Way

**Carme** – an open source framework to manage resources for multiple users running interactive jobs on an HPC Clusters



### AI development

- develop code
- schedule & track experiments
- debug model performance

### AI education

- hands-on sessions
- slides & homework
- easy Q&A

<https://carmeteam.github.io/Carme/>



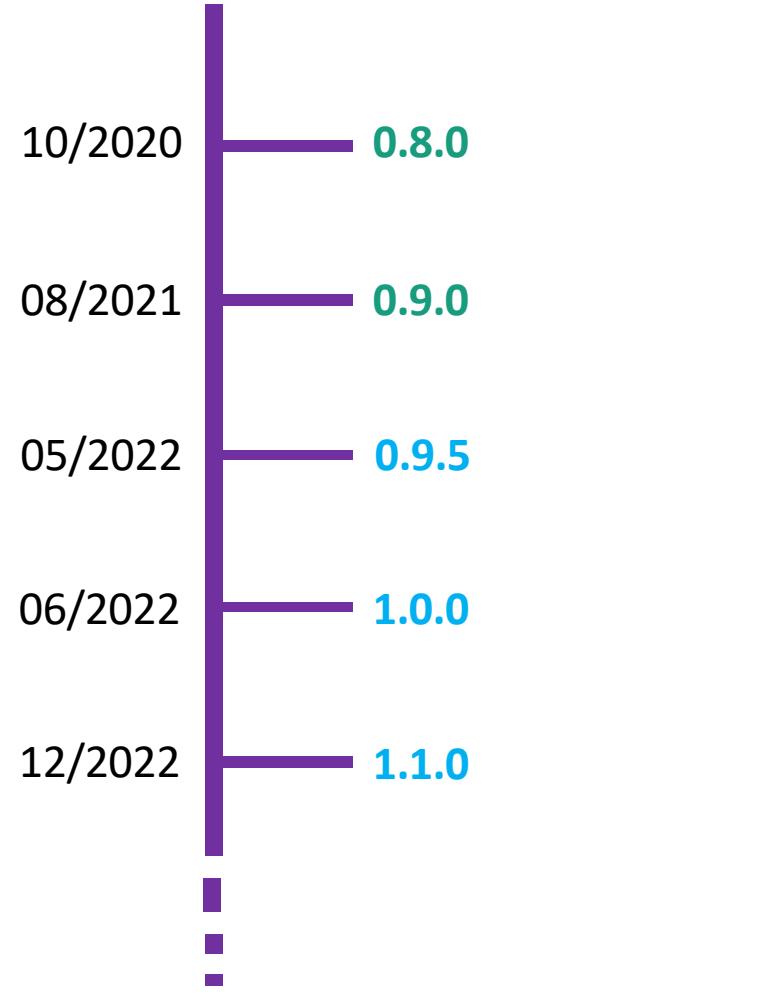
# Roadmap

—  
Where to go?

# Roadmap

## Where to go?

- available on **GitHub**  
(<https://carmeteam.github.io/Carme/>)
- **installation on demand**  
(including support)
- **release cycle:** 6–9 months
- **latest release:** r0.9.5 (05/2022)
- **next release:** r1.0.0 (06/2022)





© iStock / Fraunhofer ITWM

# Contact

---

Dr. Dominik Straßel  
Competence Center High Performance Computing  
Fraunhofer ITWM  
Fraunhofer–Platz 1  
D–67663 Kaiserslautern, GERMANY  
Tel. +49 631 31600 4896  
[dominik.strassel@itwm.fraunhofer.de](mailto:dominik.strassel@itwm.fraunhofer.de)